

User Guide for the SPEX Software Package

VERSION 3.3.0, May 30, 2025

Jinhao Chen, Timothy A. Davis, Christopher Lourenco, Lorena
Mejia-Domenzain, Erick Moreno-Centeno
Texas A&M University and US Naval Academy

Contact Information: Contact Chris Lourenco, chrisjlourenco@gmail.com,
lourenco@usna.edu, or Tim Davis, timdavis@aldenmath.com,
davis@tamu.edu, DrTimothyAldenDavis@gmail.com

1	SPEX Overview	6
2	Setting up SPEX	8
2.1	Licensing	8
2.2	Installation	8
2.2.1	Compiling and Installing SPEX using <code>cmake</code>	8
2.2.2	Compiling and Installing SPEX using <code>make</code>	9
2.2.3	Installation location	9
2.2.4	Using SPEX in your C/C++ program	9
2.2.5	Using SPEX in MATLAB and Python	9
2.2.6	Configuration options	10
3	General SPEX Data Structures and Macros	11
3.1	<code>SPEX_VERSION</code> : the version and date of the SPEX package	11
3.2	<code>SPEX_info</code> : status codes returned by SPEX	12
3.3	<code>SPEX_options</code> structure	12
3.3.1	<code>SPEX_pivot</code> : enum for pivoting schemes	13
3.3.2	<code>SPEX_preorder</code> : enum to select ordering option	14
3.3.3	<code>SPEX_factorization_algorithm</code> : enum to select factorization method	14
3.4	<code>SPEX_vector</code> : a sparse vector for updatable sparse matrices	15
3.5	The <code>SPEX_matrix</code> structure	16
3.5.1	<code>SPEX_kind</code> : enum for matrix formats	17
3.5.2	<code>SPEX_type</code> : enum for data types of matrix entries	17
3.5.3	<code>SPEX_matrix</code> structure	17
3.6	The <code>SPEX_symbolic_analysis</code> structure	20
3.6.1	<code>SPEX_factorization_kind</code> : enum for kind of factorization	20
3.6.2	<code>SPEX_symbolic_analysis</code> : containing a symbolic factorization	20
3.7	<code>SPEX_factorization</code> : containing a numeric factorization	21
4	SPEX Utilities Module	24
4.1	SPEX Utilities Module Overview	24
4.2	Managing the SPEX environment	24
4.2.1	<code>SPEX_initialize</code> : initialize the working environment	24
4.2.2	<code>SPEX_initialize_expert</code> : initialize environment (expert version)	25
4.2.3	<code>SPEX_finalize</code> : free the working environment	25
4.2.4	<code>SPEX_thread_initialize</code> : initialize working environment for a single thread	25
4.2.5	<code>SPEX_thread_finalize</code> : finalize the working environment for a single thread	26
4.3	Memory Management	26
4.3.1	<code>SPEX_calloc</code> : allocate initialized memory	26
4.3.2	<code>SPEX_malloc</code> : allocate uninitialized memory	26
4.3.3	<code>SPEX_realloc</code> : resize allocated memory	27
4.3.4	<code>SPEX_free</code> : free allocated memory	27
4.4	<code>SPEX_options</code> helper function	28

4.4.1	SPEX_create_default_options: create default SPEX_options structure	28
4.5	SPEX_vector helper functions	28
4.5.1	SPEX_vector_allocate: Allocate a SPEX vector	28
4.5.2	SPEX_vector_realloc: Reallocate a SPEX vector	29
4.5.3	SPEX_vector_free: Free a SPEX vector	29
4.6	SPEX_matrix helper functions	29
4.6.1	SPEX_matrix_allocate: allocate an m-by-n SPEX_matrix	29
4.6.2	SPEX_matrix_free: free a SPEX_matrix	30
4.6.3	SPEX_matrix_copy: make a copy of a SPEX_matrix with a potentially different matrix-format and data-type	30
4.6.4	SPEX_matrix_nnz: get the number of entries in a SPEX_matrix	31
4.6.5	SPEX_matrix_check: check and optionally print a SPEX_matrix	31
4.7	SPEX_symbolic_analysis helper function	32
4.7.1	SPEX_symbolic_analysis_free: free a symbolic analysis structure	32
4.8	SPEX_factorization helper functions	32
4.8.1	SPEX_factorization_check: check correctness of a factorization structure	32
4.8.2	SPEX_factorization_convert: Convert between updatable and non-updatable	33
4.8.3	SPEX_factorization_free: Free a SPEX factorization	33
4.9	Utility Functions	33
4.9.1	SPEX_version: Return version of the code	33
4.9.2	SPEX_determine_symmetry: Determine if a matrix is symmetric	34
4.9.3	SPEX_transpose: Transpose a CSC mpz matrix	34
4.10	SPEX_gmp: SPEX wrapper functions for GMP/MPFR	35
4.11	SPEX Helper Macros	38
4.11.1	SPEX_TRY and SPEX_CATCH	38
4.11.2	SPEX_1D: Access matrix entries with 1D linear indexing	39
4.11.3	SPEX_2D: Access dense matrix entries with 2D indexing	39
5	SPEX LU Module	40
5.1	SPEX LU Module Overview	40
5.2	Licensing	40
5.3	LU Factorization and Solve Routines	40
5.3.1	SPEX_lu_analyze: symbolic analysis for LU factorization	40
5.3.2	SPEX_lu_factorize: Compute the LU factorization of A	41
5.3.3	SPEX_lu_solve: solve the linear system	41
5.3.4	SPEX_lu_backslash: solve a linear system using an LU factorization	42
6	SPEX Cholesky and LDL Module	43
6.1	SPEX Cholesky and LDL Module Overview	43
6.2	Licensing	43
6.3	Cholesky/LDL Factorization and Solve Routines	43
6.3.1	SPEX_cholesky_analyze: symbolic analysis for Cholesky factorization	44
6.3.2	SPEX_cholesky_factorize: Compute the Cholesky factorization of A	44

6.3.3	SPEX_cholesky_solve: solve the linear system	45
6.3.4	SPEX_cholesky_backslash: solve a linear system using a Cholesky factorization	45
6.3.5	SPEX_ldl_analyze: symbolic analysis for LDL factorization	46
6.3.6	SPEX_ldl_factorize: Compute the LDL factorization of A	46
6.3.7	SPEX_ldl_solve: solve the linear system	47
6.3.8	SPEX_ldl_backslash: solve a linear system using an LDL factorization	48
7	SPEX Backslash Module	49
7.1	SPEX Backslash Module Overview	49
7.2	Licensing	49
7.3	SPEX_backslash: exactly solve a sparse linear system	49
8	SPEX Update Module	51
8.1	SPEX Update Module Overview	51
8.2	Licensing	51
8.3	SPEX Factorization Update Functions	51
8.3.1	SPEX_update_cholesky_rank1: rank-1 Cholesky update/downdate	52
8.3.2	SPEX_update_lu_colrep: Column-replacement LU update	53
8.4	SPEX Update Solve Routines	53
8.4.1	SPEX_update_solve: Solve linear system with the updatable factorization	54
8.4.2	SPEX_update_tsolve: Solve transposed linear system with the updatable factorization	54
8.5	SPEX Update Helper Function	55
8.5.1	SPEX_update_matrix_colrep: Swap the columns of a matrix	55
9	Using SPEX in MATLAB	56
9.1	Optional parameter settings	56
9.2	SPEX MATLAB functions	57
9.2.1	spex_lu_backslash.m: solve a linear system via exact LU factorization	57
9.2.2	spex_cholesky_backslash.m: solve a linear system via exact Cholesky factorization	58
9.2.3	spex_ldl_backslash.m: solve a linear system via exact LDL factorization	58
9.2.4	spex_backslash.m: solve a linear system via an exact factorization	59
9.2.5	spex_mex_demo.m: example of how to use SPEX in MATLAB	59
10	Using SPEX in Python	60
10.1	Optional parameter settings	60
10.2	SPEX Python functions	60
10.2.1	lu_backslash: solves a linear system via exact LU factorization	60
10.2.2	cholesky_backslash: solves a linear system via exact Cholesky factorization	61
10.2.3	backslash: solves a linear system via exact factorization	61

10.3 Python demo: example of how to use SPEX in Python	62
--	----

CHAPTER 1

SPEX OVERVIEW

SPEX is a software package comprising several state-of-the-art SParse EXact linear algebra routines for solving sparse linear systems exactly. One of the key goals of SPEX LU is to provide a framework for other solvers to benchmark the reliability and stability of their linear solvers, as our final solution vector x is guaranteed to be exact. SPEX LU is written in ANSI C and is accompanied by a MATLAB interface. It currently consists of the following Modules, each in their own subfolder of the SPEX package:

FIXME: I have renamed the sub-packages as Modules, to match CHOLMOD. FIXME: allow each Module be selected / unselected, like CHOLMOD.

- **SPEX_Uutilities**: Utility and auxiliary functions for all SPEX routines: interface to the GMP and MPFR libraries, memory management functions, the `SPEX_matrix`, `SPEX_factorization`, and `SPEX_symbolic_analysis` data structures, and various functions that are auxiliary to the factorization and solve functions. Refer to Chapter 4 for further details.
- **SPEX_LU**: Sparse exact left-looking LU factorization to solve the linear system $A\mathbf{x} = \mathbf{b}$. The solution time is proportional to the arithmetic work in the bit-complexity model; which is asymptotically efficient. Refer to Chapter 5 for further details.
- **SPEX_Cholesky**: Sparse exact left-looking and up-looking factorizations to solve the symmetric with nonzero leading principle minors linear system $A\mathbf{x} = \mathbf{b}$. The solution time is proportional to the arithmetic work in the bit-complexity model; this is an asymptotically efficient complexity bound. The methods can perform either a Cholesky or LDL factorization depending on the signs of the diagonal elements. Refer to Chapter 6 for further details.
- **SPEX_Backslash**: Routines to exactly solve the system $A\mathbf{x} = \mathbf{b}$ using either LU or LDL factorization. This is the simplest way to access the various Modules of SPEX. Refer to Chapter 7 for further details.
- **SPEX_LU_ColRep**: Sparse exact factorization update of an LU factorization with column replacement. Refer to Chapter 8 for further details.
- **SPEX_Cholesky_Rank1**: Sparse exact factorization update of a Cholesky or LDL factorization after a rank-1 modification (update/downdate). Refer to Chapter 8 for further details.

- **SPEX_Update_Uutilities**: Supporting methods for the sparse exact factorization update/downdate. Refer to Chapter 8 for further details.

Location: <https://github.com/clouren/SPEX> and www.suitesparse.com

Required Packages: SPEX depends on the following packages:

- GNU GMP [6] and MPFR [5] libraries. Distributed under the LGPL3 and GPL2 and can be acquired and installed from <https://gmplib.org/> and <http://www.mpfr.org/>, respectively.
- CMake, available under a BSD 3-clause license. May be independently obtained at <https://cmake.org>.
- AMD [1, 2], available under a BSD 3-clause license and distributed along with SPEX. May be independently obtained at www.suitesparse.com
- COLAMD [4, 3], available under a BSD 3-clause license and distributed along with SPEX. May be independently obtained at www.suitesparse.com
- SuiteSparse_config, available under a BSD 3-clause license and distributed along with SPEX. May be independently obtained at www.suitesparse.com.

CHAPTER 2

SETTING UP SPEX

2.1 Licensing

Copyright: The copyright of this software is held by Christopher Lourenco, Jinhao Chen, Lorena Mejia-Domenzain, Erick Moreno-Centeno, and Timothy A. Davis.

Contact Info: Chris Lourenco, chrisjlourenco@gmail.com lourenco@usna.edu, or Tim Davis, timdavis@aldenmath.com, DrTimothyAldenDavis@gmail.com, or davis@tamu.edu

License: SPEX is dual licensed under the GNU General Public License version 2 or the GNU Lesser General Public License version 3. Details of this license are in `SPEX/License/license.txt`. For alternative licenses, contact the authors.

2.2 Installation

You must first install the GMP (v6.1.2 or later), MPFR (v4.0.2 or later) and CMake (v3.22 or later) packages. See <https://gmplib.org/>, <https://www.mpfr.org/>, and <https://cmake.org/> for details. The libraries may also be installed with package managers such as spack (<https://spack.io/>) or homebrew (<https://brew.sh/>).

2.2.1 Compiling and Installing SPEX using cmake

Installation of SPEX requires the `cmake` utility in Linux, MacOS, and Windows. In a terminal window (a Command Window in Windows) go to the top-level folder of the source distribution, where you will see the subfolders called `SPEX`, `AMD`, `COLAMD`, and `SuiteSparse_config` side-by-side, and then type the following commands. This will configure the SPEX Modules, compile them, and install them as a single SPEX binary library:

```
cd build
cmake ..
cmake --build . --config Release
cmake --install .
```

The last command may require you to authenticate, for permission to install `SPEX` and dependents (`AMD`, `COLAMD`, and `SuiteSparse_config`). On Linux or Mac, use this instead of last command above:


```
sudo cmake --install .
```

2.2.2 Compiling and Installing SPEX using make

If you have the `make` command, you can instead type these commands, in the same folder, which call upon `cmake` to do the same thing as above:

```
make
sudo make install
```

Then to run some demo programs, type `make demos`.

2.2.3 Installation location

If you cannot install SPEX in the default location, you can use `cmake` to change the installation location. Instead of the command `cmake ..` above, use `ccmake ..` which will pop up a list of options. Navigate down the list and select a different location for the `CMAKE_INSTALL_PREFIX` option.

If using `make`, try `make local` instead of `make`, which will configure SPEX and its dependent libraries in the `lib` folder located within the top-level folder. To compile and install SPEX system-wide for all users on your system, use `make global ; sudo make install`.

2.2.4 Using SPEX in your C/C++ program

After compiling and installing SPEX, you can use the code inside of a C/C++ program by including `#include "SPEX.h"`, and linking against SPEX and its dependencies. CMake `find_package` scripts or `pkgconfig` files are provided as well, to use in your own cmake scripts to find SPEX and its dependencies. They are located in the `lib/cmake` and `lib/pkgconfig` folders.

2.2.5 Using SPEX in MATLAB and Python

SPEX also includes MATLAB and Python interfaces. To install the MATLAB interface, you must first configure the SPEX library outside of MATLAB and build its dependencies. Follow the instructions above for configuring, compiling, and installing SPEX (installation is optional in this case).

Next, navigate to the `SPEX/MATLAB` folder from the MATLAB Command Window and type `spex_mex_install` which will install the MATLAB interfaces to all SPEX Modules. These MATLAB functions can then be used outside of the `SPEX/MATLAB` folder by using the MATLAB `addpath` command, or `pathtool`.

The Python interface does not need any additional installation, but does require the `Numpy`, `SciPy`, and `ctypes` libraries.

2.2.6 Configuration options

Cmake can be configured with options that control the creation of Python interface for SPEX, and the use of OpenMP. OpenMP is used in SPEX for thread-local storage. If OpenMP is not available, the compiler-supported thread-local storage mechanism is used instead (if available). Each of these options can be controlled just for SPEX, or for all packages in SuiteSparse:

- **SUITESPARSE_USE_PYTHON:**

If **ON**, build all Python interfaces for SuiteSparse packages (currently only in SPEX).
If **OFF**: do not build any Python interfaces. Default: **ON**.

- **SUITESPARSE_USE_OPENMP:**

If **ON**, OpenMP is used in SuiteSparse if it is available. Default: **ON**.

- **SPEX_USE_PYTHON:**

If **ON**, build Python interface for SPEX. If **OFF**: do not build the SPEX Python interface.
Default: **SUITESPARSE_USE_PYTHON**.

- **SPEX_USE_OPENMP:**

If **ON**, OpenMP is used in SPEX if it is available. Default: **SUITESPARSE_USE_OPENMP**.

CHAPTER 3

GENERAL SPEX DATA STRUCTURES AND MACROS

The following macros/data structures are defined in `SPEX.h` and are used in all SPEX functions.

3.1 SPEX_VERSION: the version and date of the SPEX package

SPEX defines the following macros with `#define`. Refer to the `SPEX.h` file for details.

Macro	Purpose	Type
<code>SPEX_DATE</code>	Release date	string
<code>SPEX_VERSION_STRING</code>	Current version of the code	string
<code>SPEX_VERSION_MAJOR</code>	Major version of the code	integer
<code>SPEX_VERSION_MINOR</code>	Minor version of the code	integer
<code>SPEX_VERSION_SUB</code>	Sub version of the code	integer
<code>SPEX_VERSION</code>	Current version of the code	integer
<code>SPEX__VERSION</code>	Identical to <code>SPEX_VERSION</code>	integer
<code>SPEX_VERSION_NUMBER</code>	Macro to create <code>SPEX_VERSION</code>	macro

`SPEX_VERSION_NUMBER(major,minor,sub)` is a macro that creates a single integer from the three integers `major`, `minor`, `sub`. It is useful for checking relative version numbers. For example, a user application could include the following:

```
#include "SPEX.h"
#if SPEX_VERSION < SPEX_VERSION_NUMBER(3,2,0)
// SPEX version is prior to 3.2.0
#else
// SPEX is version 3.2.0 or later
#endif
```

`SPEX_VERSION` and `SPEX__VERSION` are identical. The latter is included because it follows the form of `Package__VERSION` for all packages in SuiteSparse.

See also the `SPEX_version` function described in [Section 4.9.1](#).

3.2 SPEX_info: status codes returned by SPEX

Nearly all SPEX functions return their status to the caller as their return value, an enumerated type called `SPEX_info`. The only exceptions to this rule are the memory management functions, `SPEX_malloc` and related methods, described in Section 4.3. The values of `SPEX_info` are listed below.

0	<code>SPEX_OK</code>	The function was successfully executed.
-1	<code>SPEX_OUT_OF_MEMORY</code>	Out of memory
-2	<code>SPEX_SINGULAR</code>	The input matrix A is exactly singular.
-3	<code>SPEX_INCORRECT_INPUT</code>	One or more input arguments are incorrect.
-4	<code>SPEX_NOTSPD</code>	The input matrix is not symmetric positive definite (thus can't use Cholesky)
-5	<code>SPEX_INCORRECT_ALGORITHM</code>	The algorithm is not compatible with the factorization
-6	<code>SPEX_PANIC</code>	SPEX environment error
-7	<code>SPEX_ZERODIAG</code>	Diagonal element is zero thus can't use LDL
-8	<code>SPEX_UNSYMMETRIC</code>	Input matrix is unsymmetric thus can't use LDL or Cholesky

3.3 SPEX_options structure

The `SPEX_options` struct stores key parameters for various functions used in SPEX. The `SPEX_options` option struct contains the following components:

- `option->pivot`: An enum `SPEX_pivot` type which controls the type of pivoting used. Default value: `SPEX_SMALLEST` (0). See Section 3.3.1 for details.
- `option->tol`: A double tolerance for the tolerance-based pivoting scheme, i.e., `SPEX_TOL_SMALLEST` or `SPEX_TOL_LARGEST`. `option->tol` must be in the range of $(0, 1]$. Default value: 1 meaning that the diagonal entry will be selected if it has the same magnitude as the smallest entry in the k the column. This option is only used for LU factorization. See Section 3.3.1 for details.
- `option->order`: An enum `SPEX_preorder` type which controls what column ordering is used. Default value: `SPEX_DEFAULT_ORDERING` which gives COLAMD for LU and AMD for Cholesky and LDL. See Section 3.3.2 for details.
- `option->print_level`: An int which controls the amount of output: 0: print nothing, 1: just errors, 2: terse, with basic stats from COLAMD/AMD and SPEX, 3: all, with matrices and results. Default value: 0.
- `option->prec`: A `uint64_t` which specifies the precision used for multiple precision floating point numbers, (i.e., MPFR). This can be any integer larger than `MPFR_PREC_MIN` (value of 1 in MPFR 4.0.2 and 2 in some legacy versions) and smaller than `MPFR_PREC_MAX` (usually the largest possible integer available in your system). Default value: 128 (quad precision).

- **option->round:** A `mpfr_rnd_t` which determines the type of MPFR rounding to be used by SPEX. This is a parameter of the MPFR library. The options for this parameter are:
 - `MPFR_RNDN`: Round to nearest (roundTiesToEven in IEEE 754-2008)
 - `MPFR_RNDZ`: Round toward zero (roundTowardZero in IEEE 754-2008)
 - `MPFR_RNDU`: Round toward plus infinity (roundTowardPositive in IEEE 754-2008)
 - `MPFR_RNDD`: Round toward minus infinity (roundTowardNegative in IEEE 754-2008)
 - `MPFR_RNDA`: Round away from zero
 - `MPFR_RNDF`: Faithful rounding. This is not stable.

Refer to the MPFR User Guide available at <https://www.mpfr.org/mpfr-current/mpfr.pdf> for details on the MPFR rounding style and any other utilized MPFR convention. Default value: `MPFR_RNDN`.

- **option->algo:** A `SPEX_factorization_algorithm` which indicates which type of factorization is being used. See Section 3.3.3 for details.

Most SPEX routines require **option** as their last input argument. The exceptions to this rule are the basic memory management routines in Section 4.3, the initialization/finalize methods in Section 4.2, the methods to create/free the **option** struct (Section 4.4), and the `SPEX_mp*` and `SPEX_gmp_*` wrapper functions for GMP and MPFR (Section 4.10).

The construction of the **option** struct can be avoided by passing `NULL` for the default settings, in which case default options are used. Otherwise, the following functions create and destroy a `SPEX_options` structure:

Function/Macro Name	Description	Section
<code>SPEX_create_default_options</code>	create and return <code>SPEX_options</code> pointer with default parameters upon successful allocation	4.4.1
<code>SPEX_FREE</code>	destroy <code>SPEX_options</code> structure	4.3.4

3.3.1 `SPEX_pivot`: enum for pivoting schemes

There are six available pivoting schemes provided in SPEX that can be selected with the `SPEX_options` structure. Currently, these pivot options are only used by the LU factorization, as the symmetric routines pivot exclusively down the diagonal in order to maintain symmetry. If the matrix is non-singular (in an exact sense), then the pivot is always nonzero, and is chosen as the *smallest* nonzero entry, with the smallest magnitude. This may seem counter-intuitive, but selecting a small nonzero pivot leads to smaller growth in the number of digits in the entries of `L` and `U`. This choice does not lead to any kind of numerical inaccuracy, since SPEX is guaranteed to find an exact roundoff-error free factorization of a non-singular matrix (unless it runs out of memory), for any nonzero pivot choice.

The pivot tolerance for two of the pivoting schemes is specified by the `tol` component in `SPEX_options`. The pivoting schemes for the `options->pivot` parameter are as follows:

0	SPEX_SMALLEST	The k -th pivot is selected as the smallest entry in the k -th column. This is the default.
1	SPEX_DIAGONAL	The k -th pivot is selected as the diagonal entry. If the diagonal entry is zero, this method instead selects the smallest pivot in the column.
2	SPEX_FIRST_NONZERO	The k -th pivot is selected as the first eligible nonzero in the column.
3	SPEX_TOL_SMALLEST	The k -th pivot is selected as the diagonal entry if the diagonal is within a specified tolerance of the smallest entry in the column. Otherwise, the smallest entry in the k -th column is selected. This is the default pivot selection strategy.
4	SPEX_TOL_LARGEST	The k -th pivot is selected as the diagonal entry if the diagonal is within a specified tolerance of the largest entry in the column. Otherwise, the largest entry in the k -th column is selected.
5	SPEX_LARGEST	The k -th pivot is selected as the largest entry in the k -th column.

3.3.2 SPEX_preorder: enum to select ordering option

The SPEX Library provides three ordering schemes: no ordering, COLAMD, and AMD. In LU factorization, the ordering is applied only to the columns, that is this ordering gives the matrix Q . In Cholesky and LDL factorizations, the ordering is applied to both the rows and columns, that is the ordering gives the matrices P and Q . The `option->order` parameter can be one of:

0	SPEX_DEFAULT_ORDERING	Use COLAMD for LU factorization. Use AMD for Cholesky or LDL factorization.
1	SPEX_NO_ORDERING	No pre-ordering is performed on the matrix A , that is $Q = I$.
2	SPEX_COLAMD	The permutation Q is found with COLAMD. [3]. This is recommended for LU factorization.
3	SPEX_AMD	The permutation Q is found with AMD [2]. This is recommended for Cholesky and LDL factorization.

For LU factorization, PAQ is factorized where P is determined during the numerical factorization using the method specified by `SPEX_pivot` (Section 3.3.1). For Cholesky or LDL factorization, PAQ is factorized where $P = Q^T$ is computed during the symbolic analysis. In both cases, Q is determined solely by the symbolic analysis.

3.3.3 SPEX_factorization_algorithm: enum to select factorization method

This num tells SPEX which factorization algorithm should be used, via the `option->algo` parameter.

0	<code>SPEX_ALGORITHM_DEFAULT</code>	See below.
1	<code>SPEX_LU_LEFT</code>	Left-looking LU factorization
2	<code>SPEX_CHOL_LEFT</code>	Left-looking Cholesky factorization
3	<code>SPEX_CHOL_UP</code>	Up-looking Cholesky factorization
4	<code>SPEX_LDL_LEFT</code>	Left-looking LDL factorization
5	<code>SPEX_LDL_UP</code>	Up-looking LDL factorization

This option is used slightly differently within each family of algorithms, in `SPEX*_analyze`, `SPEX*_factorize`, and `SPEX*_backslash`:

- `SPEX_lu_*`: the option can be `SPEX_ALGORITHM_DEFAULT` or `SPEX_LU_LEFT`. In both cases, a left-looking LU factorization method is used to analyze and factorize the matrix.
- `SPEX_cholesky_*`: the option can be `SPEX_ALGORITHM_DEFAULT`, `SPEX_CHOL_LEFT`, or `SPEX_CHOL_UP`. The default is to use an up-looking Cholesky factorization. Alternatively, `SPEX_CHOL_LEFT` selects a left-looking Cholesky method.
- `SPEX_ldl_*`: the option can be `SPEX_ALGORITHM_DEFAULT`, `SPEX_LDL_LEFT`, or `SPEX_LDL_UP`. The default is to use an up-looking LDL factorization. Alternatively, `SPEX_LDL_LEFT` selects a left-looking LDL method.
- `SPEX_backslash`: any option is permitted. The default algorithm is selected with `SPEX_ALGORITHM_DEFAULT`. In this case, an up-looking LDL factorization is first tried. If the matrix is unsymmetric, or if the factorization encounters zeros on the diagonal D , then the LDL factorization fails, and a left-looking LU factorization is then used. If the option is set to a particular algorithm, then it will be used instead. If the selected algorithm fails (say LDL is requested but the matrix is unsymmetric), SPEX does not fall back to trying a different algorithm (such as an LU factorization). This fallback is only available if the option is set to `SPEX_ALGORITHM_DEFAULT`.

3.4 SPEX_vector: a sparse vector for updatable sparse matrices

`SPEX_vector` is a compressed sparse vector data structure used for SPEX dynamic CSC matrices.

It is used when calling the update/downdate functions in SPEX to modify to construct the vector to modify original matrix A (either w for $A' = A + \sigma w w^T$ in rank-1 update/downdate or vk to be swapped with $A \rightarrow v[k]$ in the update for column replacement).

This is **NOT** intended to be used for building any n -by-1 vector (e.g., the right-hand-side vector b in $Ax = b$), which should be considered as a n -by-1 `SPEX_matrix`.

The `SPEX_vector` struct contains the following components:

- `vector->nz`: The number of explicit entries in the vector. Data Type: `int64_t`.

- **vector->nzmax**: The size of the **i** and **x** arrays. Note that $\text{nz} \leq \text{nzmax}$. Data Type: `int64_t`.
- **vector->i**: An array of size **nzmax** containing the row indices of all explicit entries in the vector. The last (**nzmax-nz**) entries are unused and the contents of this sparse are undefined. Data Type: `int64_t *`.
- **vector->x**: An array of size **nzmax** containing the numeric values of all explicit entries in the vector. The last (**nzmax-nz**) entries are unused and the contents of this space is undefined. Data Type: `mpz_t *`.
- **vector->scale**: Scaling parameter. The actual value of the k -th nonzero should be computed as $\text{x}[k] * \text{scale}$. Both $\text{x}[k] * \text{scale}$ and $\text{x}[k] / \text{mpq_denref}(\text{scale})$ must be integer for all entries, where `mpq_denref(scale)` is a GMP macro that gives the denominator of **scale**. This is used to skip explicit update(s) for a column/row of the factorization matrix, when all entries are to be multiplied with the same scaling factor(s). Data Type: `mpq_t`.

SPEX has a set of functions to allocate, destroy and reallocate a SPEX vector, `SPEX_vector`, as shown in the following table.

Function Name	Description	Section
<code>SPEX_vector_allocate</code>	allocate <code>SPEX_vector</code> with nzmax entries	4.5.1
<code>SPEX_vector_realloc</code>	reallocate <code>SPEX_vector</code> with new_size entries	4.5.2
<code>SPEX_vector_free</code>	destroy a <code>SPEX_vector</code> and free its allocated memory	4.5.3

3.5 The SPEX_matrix structure

SPEX operates on matrices stored in any of the 16 different matrix formats: 15 of which are combinations of matrix formats and entry data-types: $\{\text{Static Compressed Sparse Column (CSC), triplet, dense}\} \times \{\text{mpz_t, mpq_t, mpfr_t, int64_t, or double}\}$, and the 16th of which is the dynamic CSC matrix with `mpz_t` entries. Using the SPEX matrix copy function, a matrix of any given form and data-type can be copied and converted into a matrix of any one of the 16 matrix-form and data-type combinations.

Most routines require the matrix to be in CSC form with `mpz_t` (i.e., arbitrary-sized integer) data type. This data structure stores the matrix A as a sequence of three arrays:

- **A->p**: Column pointers; an array of size **n+1**. The row indices of column j are located in positions **A->p[j]** to **A->p[j+1]-1** of the array **A->i**. Data type: `int64_t`.
- **A->i**: Row indices; an array of size equal to the number of entries in the matrix. The entry **A->i[k]** is the row index of the k th nonzero in the matrix. Data type: `int64_t`.
- **A->x**: Numeric entries. The entry **A->x[k]** is the numeric value of the k th nonzero in the matrix. The array **A->x** has a union type and must be accessed via a suffix according to the type of **A**. For details, refer to Section [3.5](#).

An example matrix A with `mpz_t` type is stored as follows (note that indexing is zero based as per the C convention).

$$A = \begin{bmatrix} 1 & 0 & 0 & 1 \\ 2 & 0 & 4 & 12 \\ 7 & 1 & 1 & 1 \\ 0 & 2 & 3 & 0 \end{bmatrix}$$

```
A->p      = [0, 3, 5, 8, 11]
A->i      = [0, 1, 2, 2, 3, 1, 2, 3, 0, 1, 2]
A->x.mpz = [1, 2, 7, 1, 2, 4, 1, 3, 1, 12, 1]
```

For example, the last column appears in positions 8 to 10 of `A->i` and `A->x.mpz`, with row indices 0, 1, and 2, and values $a_{03} = 1$, $a_{13} = 12$, and $a_{23} = 1$.

3.5.1 SPEX_kind: enum for matrix formats

The SPEX library provides four available matrix formats: sparse CSC (compressed sparse column), sparse triplet, dense and sparse dynamic CSC.

0	SPEX_CSC	Matrix is in compressed sparse column format.
1	SPEX_TRIPLET	Matrix is in sparse triplet format.
2	SPEX_DENSE	Matrix is in dense format.
3	SPEX_DYNAMIC_CSC	Matrix is in dynamic CSC format.

3.5.2 SPEX_type: enum for data types of matrix entries

The SPEX library provides five data types for matrix entries: `mpz_t`, `mpq_t`, `mpfr_t`, `int64_t` and `double`.

0	SPEX_MPZ	Matrix entries are in <code>mpz_t</code> type: an integer of arbitrary size.
1	SPEX_MPQ	Matrix entries are in <code>mpq_t</code> type: a rational number with arbitrary-sized integer numerator and denominator.
2	SPEX_MPFR	Matrix entries are in <code>mpfr_t</code> type: a floating-point number of arbitrary precision.
3	SPEX_INT64	Matrix entries are in <code>int64_t</code> type.
4	SPEX_FP64	Matrix entries are in <code>double</code> type.

3.5.3 SPEX_matrix structure

A matrix `SPEX_matrix A` has the following components:

- `A->kind`: Indicating the kind/format of matrix `A`: CSC, triplet, dense or dynamic CSC. Data Type: `SPEX_kind`.
- `A->type`: Indicating the type of entries in matrix `A`: `mpz_t`, `mpq_t`, `mpfr_t`, `int64_t` or `double`. Data Type: `SPEX_type`.

- **A->m**: Number of rows in the matrix. Data Type: `int64_t`.
- **A->n**: Number of columns in the matrix. Data Type: `int64_t`.
- **A->scale**: A scaling parameter for matrix of `mpz_t` type. For all matrices whose entries are stored in data type other than `mpz_t`, SPEX assumes and maintains `A->scale = 1`. This is used to ensure that entry can be represented as an integer in an `mpz_t` matrix if these entries are converted from non-integer type data (such as double, variable precision floating point, or rational). Data Type: `mpq_t`.
- **A->nzmax**: The allocated size of the vectors `A->i`, `A->j` and `A->x`. Note that `A->nzmax` $\geq \text{nnz}(A)$, where `nnz(A)` is the value of `nnz` returned by `SPEX_matrix_nnz(&nnz,A,option)`. Data Type: `int64_t`.
- **A->nz**: The number of nonzeros in the matrix *A*, if *A* is a triplet matrix (ignored for matrices in CSC, dense or dynamic CSC formats). Data Type: `int64_t`.
- **A->p**: An array of size `A->n+1` which contains column pointers of *A*, if *A* is a CSC matrix (NULL for matrices in triplet or dense formats). Data Type: `int64_t*`.
- **A->p_shallow**: A boolean indicating whether `A->p` is shallow. A *shallow* pointer is one that refers to a component of another matrix or data structure. If `A->p` is shallow, then it should not be modified as part of the *A* matrix, and it is not freed if *A* is freed. Data Type: `bool`.
- **A->i**: An array of size `A->nzmax` which contains the row indices of the nonzeros in *A*, if *A* is a CSC or triplet matrix (NULL for dense matrices). The matrix is zero-based, so row indices are in the range of $[0, A->m-1]$. Data Type: `int64_t*`.
- **A->i_shallow**: A boolean indicating whether `A->i` is shallow. Data Type: `bool`.
- **A->j**: An array of size `A->nzmax` which contains the column indices of the nonzeros in *A*, if *A* is a triplet matrix (NULL for matrices in CSC or dense formats). The matrix is zero-based, so column indices are in the range of $[0, A->n-1]$. Data Type: `int64_t*`.
- **A->j_shallow**: A boolean indicating whether `A->j` is shallow. Data Type: `bool`.
- **A->x**: An array of size `A->nzmax` which contains the numeric values of the matrix. This array is a union, and must be accessed via one of: `A->x.mpz`, `A->x.mpq`, `A->x.mpfr`, `A->x.int64`, or `A->x.fp64`, depending on the `A->type` parameter. Data Type: `union`.
- **A->x_shallow**: A boolean indicating whether `A->x` is shallow. Data Type: `bool`.
- **A->v**: If the matrix is a `SPEX_DYNAMIC_CSC` this is an array of size `A->n`, each of which is a dynamic column vector. Data Type: `SPEX_vector**`. Always NULL in the current release of SPEX.

Specifically, for different kinds of *A* of size `A->m` $\times$ `A->n` with `nz` nonzero entries, its components are defined as:

- (0) **SPEX_CSC**: A sparse matrix in CSC (compressed sparse column) format. **A->p** is an `int64_t` array of size **A->n+1**, **A->i** is an `int64_t` array of size **A->nzmax** (with $nz \leq A \rightarrow nzmax$), and **A->x.TYPE** is an array of size **A->nzmax** of matrix entries (**TYPE** is one of `mpz`, `mpq`, `mpfr`, `int64`, or `fp64`). The row indices of column j appear in **A->i** [**A->p** [j] ... **A->p**[$j+1$]-1], and the values appear in the same locations in **A->x.TYPE**. The **A->j** array is **NULL**. **A->nz** is ignored; the number of entries in **A** is given by **A->p** [**A->n**]. Row indices need not be sorted in each column, but duplicates cannot appear.
- (1) **SPEX_TRIPLET**: A sparse matrix in triplet format. **A->i** and **A->j** are both `int64_t` arrays of size **A->nzmax**, and **A->x.TYPE** is an array of values of the same size. The k th tuple has row index **A->i** [k], column index **A->j** [k], and value **A->x.TYPE** [k], with $0 \leq k < A \rightarrow nz$. The **A->p** array is **NULL**. Triplets can be unsorted, but duplicates cannot appear.
- (2) **SPEX_DENSE**: A dense matrix. The integer arrays **A->p**, **A->i**, and **A->j** are all **NULL**. **A->x.TYPE** is a pointer to an array of size **A->m-by-A->n**, stored in column-oriented format. The value of $A(i, j)$ is **A->x.TYPE** [p] with $p = i + j * A \rightarrow m$. **A->nz** is ignored; the number of entries in **A** is **A->m** $\times$ **A->n**.
- (3) **SPEX_DYNAMIC_CSC**: A sparse matrix in dynamic CSC format with the number of nonzeros in each column changing independently and dynamically, which is only used in the **SPEX** update functions. The matrix is held as an array of **A->n** **SPEX_vector**s, one per column. Each column is held as a **SPEX_vector**, containing `mpz_t` values and its own scale factor. For this kind, **A->nzmax**, **A->nz**, **A->p**, **A->i**, **A->x** and **A->*_shallow** are ignored and pointers **p**, **i** and **x** are remained as **NULL** pointers. To access entries in column j , **A->v[j]->i**[0 ... **A->v[j]->nz-1**] give the row indices of all nonzeros, and the `mpz_t` values of these entries appear in the same locations in **A->v[j]->x**. **A->v[j]->nzmax** is the max number of nonzeros allocated for the j th column, while **A->v[j]->nz** is the number of existing nonzeros in the j th column. Therefore, the total number of existing nonzeros is computed as $\sum_{j=0}^{n-1} (A \rightarrow v[j] \rightarrow nz)$.

A may contain shallow components, **A->p**, **A->i**, **A->j**, and **A->x**. For example, if **A->p_shallow** is true, then a non-**NULL** **A->p** is a pointer to a read-only array, and the **A->p** array is not freed by **SPEX_matrix_free**. If **A->p** is **NULL** (for a triplet or dense matrix), then **A->p_shallow** has no effect.

SPEX has a set of functions to allocate, copy/convert, query and destroy a **SPEX_matrix** as shown in the following table.

Function Name	Description	Section
SPEX_matrix_allocate	allocate a m -by- n SPEX_matrix	4.6.1
SPEX_matrix_free	destroy a SPEX_matrix and free its allocated memory	4.6.2
SPEX_matrix_copy	make a copy of a matrix, into another kind and/or type	4.6.3
SPEX_matrix_nnz	get the number of entries in a matrix	4.6.4
SPEX_matrix_check	check the validity of a matrix data structure, and print it	4.6.5

3.6 The SPEX_symbolic_analysis structure

The symbolic analysis structure handles all preorderings and structural information for each factorization within SPEX. Section 3.6.1 discusses an enum for the type of factorization. Section 3.6.2 discusses the components of this data structure.

3.6.1 SPEX_factorization_kind: enum for kind of factorization

The SPEX library currently provides three types of factorizations: LU, Cholesky, and LDL. The value `SPEX_QR_FACTORIZATION` is reserved for future development.

0	<code>SPEX_LU_FACTORIZATION</code>	LU factorization is being used
1	<code>SPEX_CHOLESKY_FACTORIZATION</code>	Cholesky factorization is being used
2	<code>SPEX_LDL_FACTORIZATION</code>	LDL factorization is used
3	<code>SPEX_QR_FACTORIZATION</code>	QR factorization is being used (reserved for future use)

3.6.2 SPEX_symbolic_analysis: containing a symbolic factorization

A symbolic analysis object `SPEX_symbolic_analysis S` has the following components:

- `S->kind`: Indicating the kind of factorization either LU, Cholesky, or LDL. Data type: `SPEX_factorization_kind`.
- `S->P_perm`: Row permutation for Cholesky/LDL and LU factorization. Data type: `int64_t *`.
- `S->Pinv_perm`: Inverse row permutation for Cholesky/LDL and LU factorization. Data type: `int64_t *`.
- `S->Q_perm`: Column permutation for LU factorization. This is always `NULL` and ignored for Cholesky/LDL factorization since its row and column permutations are the same. Data type: `int64_t *`.
- `S->Qinv_perm`: Inverse column permutation for LU factorization. This is always `NULL` and ignored for Cholesky/LDL factorization since its inverse row and column permutations are the same. Data type: `int64_t *`.
- `S->lnoz`: Approximate number of nonzeros in L . In LU factorization, this is a crude estimate based on either AMD or COLAMD. In Cholesky/LDL factorization, if AMD is used, this is the exact number of nonzeros in L (excluding numeric cancellation). Data type: `int64_t`.
- `S->unoz`: Approximate number of nonzeros in U . In LU factorization, this is a crude estimate based on either AMD or COLAMD. In Cholesky/LDL factorization this is not used. Data type: `int64_t`.
- `S->parent`: This is the elimination tree of the input matrix for Cholesky or LDL factorization. This is always `NULL` for LU factorization. Data type: `int64_t *`.

- **S->cp**: Column pointers of L for Cholesky or LDL factorization. This is always NULL for LU factorization. Data type: `int64_t *`.

This data type is constructed when analysis is called in the appropriate factorizations. See sections 5.3.1 and 6.3.1 for further details. To free this data structure, the function `SPEX_symbolic_analysis_free` is used and discussed further in section 4.7.

3.7 SPEX_factorization: containing a numeric factorization

The `SPEX_factorization` object holds an LU, Cholesky, or LDL numerical factorization. in either non-updatable or updatable form. The components of the factorization structure are accessible to the user application. However, they should only be modified by calling SPEX methods. Changing them directly can lead to undefined behavior.

FIXME: for v3.3.0: only LU is updatable, not Cholesky or LDL.

The components of a `SPEX_factorization F` are as follows:

- **F->kind**: Indicating the kind of factorization either LU, Cholesky, or LDL. Data type: `SPEX_factorization_kind`.
- **F->updatable**: a flag that indicates whether the factorization is in an updatable format. Data type: `bool`.
- **F->scale_for_A**: Scaling factor of the input matrix *A*. As discussed in section 3.5, all matrices in SPEX are integral, thus, if *A* must be scaled the scaling factor applied is stored here. Data type: `mpq_t`.
- **F->L**: The lower triangular matrix for either LU, Cholesky, or LDL factorization. Data type: `SPEX_matrix`.
- **F->U**: The upper triangular matrix for LU factorization. This is always NULL for Cholesky/LDL factorization. Data type: `SPEX_matrix`.
- **F->Q**: The matrix for (future) QR factorization. Provided here so that future versions of SPEX have backward compatibility with this version of SPEX. Data type: `SPEX_matrix`.
- **F->R**: The right triangular matrix for (future) QR factorization. Provided here so that future versions of SPEX have backward compatibility with this version of SPEX. Data type: `SPEX_matrix`.
- **F->rhos**: An $n \times 1$ dense matrix containing the pivot values used for LU or Cholesky/LDL factorization. Data type: `SPEX_matrix`.
- **F->P_perm**: Row permutation of the LU or Cholesky/LDL factors. Data type: `int64_t *`.

- **F->Pinv_perm**: Inverse row permutation of the LU or Cholesky/LDL factors. Data type: `int64_t *`.
- **F->Q_perm**: Column permutation of the LU factors. This is NULL and ignored for Cholesky/LDL factorization. Data type: `int64_t *`.
- **F->Qinv_perm**: Inverse column permutation of the LU factors. This is NULL and ignored for Cholesky/LDL factorization. Data type: `int64_t *`.
- **F->rank**: rank of matrix A ; reserved for future development in a future QR facotrization.

As mentioned in the above description, one or more of these components could be NULL for certain factorization. For example, **F->U**, **F->Q_perm** and **F->Qinv_perm** are NULL for Cholesky factorization. Moreover, **F->Qinv_perm** can be NULL for a non-updatable LU factorization with **F->updatable=false**, but it will be generated when the LU factorization converted to the updatable form.

Aside from that **Qinv_perm** will be generated for the updatable LU factorization, the only difference between non-updatable and updatable factorizations is the way that L (and U if exists) is stored. Specifically, a updatable factorization must meet the following conditions.

- Both **F->L** and **F->U** are `SPEX_DYNAMIC_CSC` matrices. Notably, **F->U** in the updatable factorization is actually the transpose of U (or equivalently, U is stored as compressed sparse row (CSR) form instead), since U will be updated one row at a time.
- $A = LDU$, which means L and U are properly permuted. (Recall that $PAQ = LDU$ or $PAP^T = LDL^T$ holds for static factorization). That is, for a updatable factorization, the rows of L are in the same order as the rows of A , while the j -th column of L (**F->L->v[j]**) contains the j -th pivot, which would be **F->L->v[j]->x[0]**, (i.e., **F->L->v[j]->i[0] == F->P_perm[j]**); the columns of U (i.e., the rows of U^T) are in the same order as the columns of A , while the j -th row of U (i.e., the j -th column of U^T) (**F->U->v[j]**) contains the j -th pivot, which would be **F->U->v[j]->x[0]**, (i.e., **F->U->v[j]->i[0] == F->Q_perm[j]**).

Due to these non-trivial conditions, users cannot simply perform `SPEX_matrix_copy` to obtain **F->L** and/or **F->U** in `SPEX_DYNAMIC_CSC` `SPEX_MPZ` format and claim that it is updatable, and vice versa. To correctly convert the factorization, the user application should call `SPEX_factorization_convert` to perform in-place conversion for a given **F** to either updatable or non-updatable as specified.

A `SPEX_factorization` is constructed by the appropriate factorization algorithms (see Sections 5.3.2 and 6.3.2 for further details). To free this data structure, the function `SPEX_factorization_free` is used and discussed further in section 4.8.3., and updated by appropriate `SPEX_update_*` functions (see sections 8.3.1 and 8.3.2 for further details). In addition, SPEX has a set of functions to convert, query and destroy a SPEX factorization, `SPEX_factorization`, as shown in the following table.

Function Name	Description	Section
<code>SPEX_factorization_check</code>	check the validity of a factorization and print it	4.8.1
<code>SPEX_factorization_convert</code>	convert a factorization to non-/updatable as specified	4.8.2
<code>SPEX_factorization_free</code>	destroy a <code>SPEX_factorization</code> and free its allocated memory	4.8.3

It should be noted that all solvers (except `SPEX_update_(t)solve`) and functions that create factorization return non-updatable factorization with F->L (and F->U if exists) in `SPEX_CSC` `SPEX_MPZ` form. On the other hand, all `SPEX_update_*` functions require and output updatable factorization with F->L (and F->U if exists) in `SPEX_DYNAMIC_CSC` `SPEX_MPZ` form. These `SPEX_update_*` functions check the input factorization, convert it (if not updatable) automatically, and output updatable factorization.

CHAPTER 4

SPEX UTILITIES MODULE

4.1 SPEX Utilities Module Overview

SPEX Util contains utility and auxiliary functions for the SPEX factorizations. Additionally, SPEX Util provides a wrapper class for the GNU Multiple Precision Arithmetic (GMP) [6] and GNU Multiple Precision Floating Point Reliable (MPFR) [5] libraries that prevent memory leaks and improve the overall stability of these external libraries by removing `abort()` conditions.

4.2 Managing the SPEX environment

Either `SPEX_initialize` or `SPEX_initialize_expert` (but not both) must be called prior to using any other SPEX functions. Otherwise, all SPEX user-callable functions will return `SPEX_PANIC`. `SPEX_finalize` must be called as the last SPEX function. Note that if a user application is working in a multithreaded environment then only one user thread should call the `SPEX_initialize` and `SPEX_finalize` functions.

Subsequent SPEX sessions can be restarted after a call to `SPEX_finalize`, by calling either `SPEX_initialize` or `SPEX_initialize_expert` (but not both), followed by a final call to `SPEX_finalize` when finished.

4.2.1 `SPEX_initialize`: initialize the working environment

```
SPEX_info SPEX_initialize ( void ) ;
```

`SPEX_initialize` initializes the working environment for SPEX functions. SPEX utilizes a specialized memory management scheme in order to prevent potential memory failures caused by GMP and MPFR libraries. Either this function or `SPEX_initialize_expert` must be called prior to using any other function in the library. By default, `SPEX_initialize` utilizes the `Suitesparse_malloc`, `Suitesparse_calloc`, `SuiteSparse_realloc`, and `Suitesparse_free` functions, which default to using the ANSI C `malloc`, `calloc`, `realloc`, and `free` functions. `SPEX_initialize` returns `SPEX_PANIC` if SPEX has already been initialized, or `SPEX_OK` if successful.

4.2.2 SPEX_initialize_expert: initialize environment (expert version)

```
SPEX_info SPEX_initialize_expert
(
    void *(*MyMalloc) (size_t),           // user-defined malloc
    void *(*MyCalloc) (size_t, size_t),   // user-defined calloc
    void *(*MyRealloc) (void *, size_t),   // user-defined realloc
    void (*MyFree) (void *)               // user-defined free
);
```

`SPEX_initialize_expert` is the same as `SPEX_initialize` except that it allows for a redefinition of custom memory functions that are used for SPEX and GMP/ MPFR. The four inputs to this function are pointers to four functions with the same signatures as the ANSI C `malloc`, `calloc`, `realloc`, and `free` functions. That is:

```
#include <stdlib.h>
void *malloc (size_t size) ;
void *calloc (size_t nmemb, size_t size) ;
void *realloc (void *ptr, size_t size) ;
void free (void *ptr) ;
```

Returns `SPEX_PANIC` if SPEX has already been initialized, or `SPEX_OK` if successful.

4.2.3 SPEX_finalize: free the working environment

```
SPEX_info SPEX_finalize ( void ) ;
```

`SPEX_finalize` finalizes the working environment for SPEX library, and frees any internal workspace created by SPEX. It must be called as the last `SPEX_*` function called, except that a subsequent call to `SPEX_initialize*` may be used to start another SPEX session. Returns `SPEX_PANIC` if SPEX has not been initialized, or `SPEX_OK` if successful.

4.2.4 SPEX_thread_initialize: initialize working environment for a single thread

```
SPEX_info SPEX_thread_initialize ( void ) ;
```

`SPEX_thread_initialize` initializes the working environment of SPEX for a single user thread. If the user is working in a multithreaded environment, they must call this function at the beginning of each user thread. Returns `SPEX_OK` if successful or `SPEX_PANIC` if SPEX was already initialized.

This function is only required for a multithreaded user application that calls SPEX functions from threads other than the primary thread that called `SPEX_initialize`.

When the primary thread of the user application starts, it must call `SPEX_initialize`. When the user application enters a parallel region (say with OpenMP) or creates its own threads with a threading library, each user thread that calls any SPEX function must first call `SPEX_thread_initialize` when it starts, and `SPEX_thread_finalize` when it finishes.

An example usage can be found in the SPEX/Demo folder in the `spex_demo_threaded.c` main program.

4.2.5 `SPEX_thread_finalize`: finalize the working environment for a single thread

```
SPEX_info SPEX_thread_finalize ( void ) ;
```

`SPEX_thread_finalize` finalizes the working environment and frees any internal workspace created by SPEX for a single user thread. If the user is working in a multithreaded environment, they must call this function at the end of each user thread. Returns `SPEX_OK` if successful or `SPEX_PANIC` if SPEX was not initialized.

4.3 Memory Management

The routines in this section are used to allocate and free memory for the data structures used in SPEX. By default, SPEX relies on the SuiteSparse memory management functions, `SuiteSparse_malloc`, `SuiteSparse_calloc`, `SuiteSparse_realloc`, and `SuiteSparse_free`. By default, those functions rely on the ANSI C `malloc`, `calloc`, `realloc`, and `free`, but this may be changed by initializing the SPEX environment with `SPEX_initialize_expert`.

4.3.1 `SPEX_calloc`: allocate initialized memory

```
void *SPEX_calloc
(
    size_t nitems,      // number of items to allocate
    size_t size         // size of each item
) ;
```

`SPEX_calloc` allocates a block of memory for an array of `nitems` elements, each of them `size` bytes long, and initializes all its bits to zero. If any input is less than 1, it is treated as if equal to 1. If the function failed to allocate the requested block of memory, then a `NULL` pointer is returned. Returns `NULL` if the allocation fails.

4.3.2 `SPEX_malloc`: allocate uninitialized memory

```
void *SPEX_malloc
(
    size_t size         // size of memory space to allocate
) ;
```

`SPEX_malloc` allocates a block of `size` bytes of memory, returning a pointer to the beginning of the block. The content of the newly allocated block of memory is not initialized, remaining with indeterminate values. If `size` is less than 1, it is treated as if equal to 1. If the function fails to allocate the requested block of memory, then a `NULL` pointer is returned. Returns `NULL` if the allocation fails.

4.3.3 SPEX_realloc: resize allocated memory

```

void *SPEX_realloc      // pointer to reallocated block, or original block
                        // if the realloc failed
(
    int64_t nitems_new,    // new number of items in the object
    int64_t nitems_old,    // old number of items in the object
    size_t size_of_item,   // sizeof each item
    void *p,              // old object to reallocate
    bool *ok               // true if success, false on failure
) ;

```

`SPEX_realloc` is a wrapper for `realloc`. If `p` is non-NULL on input, it points to a previously allocated array of size `nitems_old × size_of_item`. The array is reallocated to be of size `nitems_new × size_of_item`. If `p` is NULL on input, then a new array of that size is allocated. On success, a pointer to the new array is returned.

If the reallocation fails, the space pointed to by `p` (if `p` is not NULL on input) is not modified, and `ok` is returned as `false` to indicate that the reallocation failed. If the size decreases or remains the same, then the method always succeeds (`ok` is returned as `true`). The return value of the function is the unchanged input value of `p`.

Typical usage: the following code fragment allocates an array of 10 `int`'s, and then increases the size of the array to 20 `int`'s. If the `SPEX_malloc` succeeds but the `SPEX_realloc` fails, then the array remains unmodified, of size 10.

```

int *p ;
p = SPEX_malloc (10 * sizeof (int)) ;
if (p == NULL) { error here ... }
printf ("p points to an array of size 10 * sizeof (int)\n") ;
bool ok ;
p = SPEX_realloc (20, 10, sizeof (int), p, &ok) ;
if (ok) printf ("p has size 20 * sizeof (int)\n") ;
else printf ("realloc failed; p still has size 10 * sizeof (int)\n") ;
SPEX_free (p) ;

```

4.3.4 SPEX_free: free allocated memory

```

void SPEX_free
(
    void *p              // pointer to memory space to free
) ;

```

`SPEX_free` frees the memory previously allocated by a call to `SPEX_calloc`, `SPEX_malloc`, or `SPEX_realloc`. Results are undefined if `p` is not NULL on input, but was not obtained from one of these three functions. If `p` is NULL on input, then no action is taken (this is not an error condition). To guard against freeing the same memory space twice, the following macro `SPEX_FREE` is provided, which calls `SPEX_free` and then sets the freed pointer to NULL.

```

#define SPEX_FREE(p)          \
{                               \
    SPEX_free (p) ;           \
    (p) = NULL ;              \
}

```

4.4 SPEX_options helper function

The `SPEX_options` structure contains numerous parameters that may be modified to change the behavior of the SPEX functions. Default values of these parameters will lead to good performance in most cases. The following helper functions are provided.

4.4.1 SPEX_create_default_options: create default SPEX_options structure

```
SPEX_info SPEX_create_default_options (SPEX_options *option_handle) ;
```

`SPEX_create_default_options` creates and returns a pointer to a `SPEX_options` struct with default parameters upon successful allocation, which are discussed in Section 3.3. To safely free the `SPEX_options` option structure, simply use `SPEX_FREE(option)`. All functions that require `SPEX_options` option as an input argument can have a `NULL` pointer passed instead. In this case, the default value of the corresponding option is used.

4.5 SPEX_vector helper functions

4.5.1 SPEX_vector_allocate: Allocate a SPEX vector

```

SPEX_info SPEX_vector_allocate
(
    SPEX_vector *v_handle,          // vector to be allocated
    const int64_t nzmax,            // number of nnz entries in v
    const SPEX_options option
) ;

```

`SPEX_vector_allocate` creates and initializes a `SPEX_vector` of size `nzmax`. On input, the `SPEX_vector *v` that `v_handle` points to is `NULL`. On output, `SPEX_vector *v` is allocated with `v->nz=0`, `v->nzmax=nzmax`, `v->scale=1`, `v->i` being an `int64_t` array of size `nzmax` and `v->x` being an array containing `nzmax` initialized `mpz_t` entries.

4.5.2 SPEX_vector_realloc: Reallocate a SPEX vector

```
SPEX_info SPEX_vector_realloc
(
    SPEX_vector v,           // the vector to be expanded
    const int64_t new_size,   // desired new size for v
    const SPEX_options option
);
```

`SPEX_vector_realloc` either expands or shrinks the given `SPEX_vector *v` of size `v->nzmax` so that its new size is `new_size`. If `new_size > v->nzmax`, all newly allocated `mpz_t` entries are initialized.

4.5.3 SPEX_vector_free: Free a SPEX vector

```
SPEX_info SPEX_vector_free
(
    SPEX_vector *v_handle,    // vector to be deleted
    const SPEX_options option
);
```

`SPEX_vector_free` frees the `SPEX_vector *v` that `v_handle` points to. On output, the memory associated with `v` is freed and `v` is set to `NULL`.

4.6 SPEX_matrix helper functions

4.6.1 SPEX_matrix_allocate: allocate an m-by-n SPEX_matrix

```
SPEX_info SPEX_matrix_allocate
(
    SPEX_matrix *A_handle,    // matrix to allocate
    SPEX_kind kind,           // CSC, triplet, dense, and dynamic CSC
    SPEX_type type,           // mpz, mpq, mpfr, int64, or double
    int64_t m,                // # of rows
    int64_t n,                // # of columns
    int64_t nzmax,            // max # of entries for CSC or triplet
                                // (ignored if A is dense)
    bool shallow,             // if true, matrix is shallow. A->p, A->i, A->j,
                                // A->x are all returned as NULL and must be set
                                // by the caller. All A->*_shallow are returned
                                // as true. Ignored if kind is dynamic_CSC.
    bool init,                // If true, and the data types are mpz, mpq, or
                                // mpfr, the entries are initialized (using the
                                // appropriate SPEX_mp*_init function). If false,
                                // the mpz, mpq, and mpfr arrays are malloced but
                                // not initialized. Utilized internally to reduce
                                // memory. Ignored if shallow is true.
    const SPEX_options option
);
```

`SPEX_matrix_allocate` allocates memory space for a m -by- n `SPEX_matrix` whose kind (CSC, triplet, dense, or dynamic CSC) and data type (`mpz`, `mpq`, `mpfr`, `int64` or `fp64`) is specified. On input, the `SPEX_matrix` that `A_handle` points to is `NULL`. On output, `A_handle` points to a `SPEX_matrix` of specified type, kind and size.

For a CSC, triplet or dense matrix, if `shallow` is true, all components (`A->p`, `A->i`, `A->j`, `A->x`) are returned as `NULL`, and their shallow flags are all true. The pointers `A->p`, `A->i`, `A->j`, and/or `A->x` can then be assigned from arrays in the calling application. If `shallow` is false, the appropriate individual arrays are allocated (via `SPEX_calloc`). The second boolean parameter `init` is used if the entries are `mpz_t`, `mpq_t`, or `mpfr_t`. Specifically, if `init` is true, the individual entries within `A->x.TYPE` are initialized using the appropriate `SPEX_mp*_init` function. Otherwise, if `init` is false, the `A->x.TYPE` array is allocated (via `SPEX_calloc`) and left that way. They are not otherwise initialized, and attempting to access the values of these uninitialized entries will lead to undefined behavior.

For a `SPEX_DYNAMIC_CSC` matrix, `type`, `shallow` and `init` are ignored (since it only allows `mpz_t` entries). Moreover, each column of the returned `SPEX_DYNAMIC_CSC` matrix will be allocated as `SPEX_vector` with no space for any entries. Additional reallocation for each column are performed as entries are added to the matrix.

4.6.2 `SPEX_matrix_free`: free a `SPEX_matrix`

```
SPEX_info SPEX_matrix_free
(
    SPEX_matrix *A_handle, // matrix to free
    const SPEX_options option
) ;
```

`SPEX_matrix_free` frees the `SPEX_matrix` `A`. Note that the input of the function is the pointer to the pointer of a `SPEX_matrix` structure. This is because this function internally sets the pointer of a `SPEX_matrix` to be `NULL` to prevent potential segmentation fault that could be caused by double call to `SPEX_matrix_free`.

4.6.3 `SPEX_matrix_copy`: make a copy of a `SPEX_matrix` with a potentially different matrix-format and data-type

```
SPEX_info SPEX_matrix_copy
(
    SPEX_matrix *C_handle, // matrix to create (never shallow)
    // inputs, not modified:
    SPEX_kind C_kind,      // C->kind: CSC, triplet, dense, or dynamic CSC
    SPEX_type C_type,      // C->type: mpz_t, mpq_t, mpfr_t, int64_t, or double
    const SPEX_matrix A,   // matrix to make a copy of (may be shallow)
    const SPEX_options option
) ;
```

`SPEX_matrix_copy` makes a deep copy of a `SPEX_matrix` `A` as a new `SPEX_matrix` `C`, which can be any of the 16 matrix formats discussed in Section 3.5. That is, the new matrix `C` can be exactly the same as `A` or any other type or kind different than `A`. On input, the

SPEX matrix that `C_handle` points to must be `NULL` and will be ignored, and `A` is a valid matrix that can be potentially shallow. On output, `C_handle` points to the matrix `C`, which is a copy of `A` of kind `kind` and type `type`.

Results are undefined for an invalid input matrix `A`. Though all matrices generated from any SPEX user-callable functions are valid, they could become invalid when user directly modifies their component(s). To check the validity of the input matrix, call `SPEX_matrix_check` (Section 4.6.5).

4.6.4 `SPEX_matrix_nnz`: get the number of entries in a SPEX matrix

```
SPEX_info SPEX_matrix_nnz      // find the # of entries in A
(
    int64_t *nnz,              // # of entries in A, -1 if A is invalid
    const SPEX_matrix A,       // matrix to query
    const SPEX_options option
) ;
```

`SPEX_matrix_nnz` returns an integer, `nnz`, which is equal to the number of entries in a SPEX matrix `A`. For details regarding how the number of entries is obtained for different kinds of matrices, refer to Section 3.5. For any invalid matrix, `nnz` is returned as -1.

4.6.5 `SPEX_matrix_check`: check and optionally print a SPEX matrix

```
SPEX_info SPEX_matrix_check    // returns a SPEX status code
(
    const SPEX_matrix A,       // matrix to check
    const SPEX_options option
) ;
```

`SPEX_matrix_check` checks the validity of a SPEX matrix `A` in any of the 16 matrix formats discussed in Section 3.5. In addition, it optionally prints the matrix and any error found with proper print level specified by `option->print_level`. Specifically, `SPEX_matrix_check` prints nothing for `print_level=0` (default); or just errors for `print_level=1`; or errors and terse output of the matrix for `print_level=2`; or errors and detailed output of the matrix for `print_level=3`. As mentioned, if default settings are desired, `option` can be input as `NULL`.

4.7 SPEX_symbolic_analysis helper function

4.7.1 SPEX_symbolic_analysis_free: free a symbolic analysis structure

```
SPEX_info SPEX_symbolic_analysis_free
(
    SPEX_symbolic_analysis *S_handle,    // symbolic object to be deleted
    const SPEX_options option
) ;
```

`SPEX_symbolic_analysis_free` frees the memory of the `SPEX_symbolic_analysis` `S` that `S_handle` points to. On output, the symbolic analysis `S` is set to `NULL`.

4.8 SPEX_factorization helper functions

4.8.1 SPEX_factorization_check: check correctness of a factorization structure

```
SPEX_info SPEX_factorization_check
(
    SPEX_factorization F,           // The factorization to check
    const SPEX_options option
) ;
```

`SPEX_factorization_check` checks if a given factorization `F` is correctly formatted. Specifically, it checks the following 5 conditions:

1. All required components of `F` are present (e.g., an array was not erroneously freed).
2. The sizes of all matrices match (e.g., `F->L` and `F->U` are square matrices of the same dimension).
3. `F->L` (and `F->U` if exists) is correctly formatted via `SPEX_matrix_check`.
4. `F->L`, `F->U`, and `F->rhos` have the correct and same pivot values. Additionally, when `F->updatable=true`, `F->L` (and `F->U` if exists) is of `SPEX_DYNAMIC_CSC`, and the i -th pivot is the first entry in the nonzero list of i -th vector of `L` (and `U` if exists), i.e., `F->L->v[i]->i[0] == F->P_perm[i]` and `F->U->v[i]->i[0] == F->Q_perm[i]`.
5. Each permutation has no repeated indices and maps from $[0..n-1]$ to $[0..n-1]$; `P_perm` and `Pinv_perm` are mutually inverse vectors, same applied to `(Q_perm, Qinv_perm)` if exists.

Similar to `SPEX_matrix_check`, `SPEX_factorization_check` also prints values of the factorization, together with any error found, with proper print level specified by `option->print_level`. Specifically, `SPEX_factorization_check` prints nothing for `print_level=0` (default); or just errors for `print_level=1`; or errors and terse output of the factorization for `print_level=2`; or errors and detailed output of the factorization for `print_level=3`. As mentioned, if default settings are desired, `option` can be input as `NULL`.

4.8.2 `SPEX_factorization_convert`: Convert between updatable and non-updatable

```
SPEX_info SPEX_factorization_convert
(
    SPEX_factorization F,          // The factorization to be converted
    bool updatable,               // if true: make F updatable
                                // if false: make non-updatable
    const SPEX_options option
);
```

`SPEX_factorization_convert` performs in-place conversion between updatable and non-updatable factorization as specified by the boolean input argument `updatable`. If `F->updatable` matches the `updatable` input parameter, this function does nothing. Otherwise, it performs the corresponding in-place conversion. In case of any error, the returned factorization should be considered as undefined.

Upon input, `F->L` and `F->U` (if it exists) must be non-shallow `SPEX_CSC` `SPEX_MPZ` matrices for non-updatable (static) factorization (i.e., `F->updatable == false`), otherwise, the input format is considered as incorrect and `SPEX_INCORRECT_INPUT` is returned. Likewise, `F->L` and `F->U` (if it exists) must be `SPEX_DYNAMIC_CSC` `SPEX_MPZ` matrices for an updatable factorization. All SPEX functions output a factorization in either of these two formats, and are non-shallow. Therefore, these input requirements can be met easily if the user application does not modify any individual component of `F` itself, except via calls to SPEX methods.

4.8.3 `SPEX_factorization_free`: Free a SPEX factorization

```
SPEX_info SPEX_factorization_free
(
    SPEX_factorization *F_handle, // numeric facotorization to be deleted
    const SPEX_options option
);
```

`SPEX_factorization_free` frees the memory of the `SPEX_factorization` that `F_handle` points to, and sets its contents, `F = *F_handle`, to `NULL`. The factorization `F` is one that was constructed by `SPEX_lu_factorize`, `SPEX_cholesky_factorize`, or `SPEX_ldl_factorize`.

4.9 Utility Functions

4.9.1 `SPEX_version`: Return version of the code

```
SPEX_info SPEX_version
(
    int version [3],           // SPEX major, minor, and sub version
    char date [128]           // date of this version
);
```

`SPEX_version` returns the library version and date. The `version` array contains the three version numbers that are available at compile-time `#define`'d values: `SPEX_VERSION_MAJOR`,

SPEX_VERSION_MINOR, and SPEX_VERSION_SUB, in that order. The `SPEX_version` function allows the user application to check which version of SPEX it has been linked with. The three `#define`'d values allow the user application to know which version of SPEX was used at compile-time, which might not be the same version that was linked later on. The `date` is the string `SPEX_DATE`, in the form "Mar 31, 2023" for example. The string is null-terminated.

The `SPEX_VERSION` macro is a single integer with the entire version information, created by the macro `SPEX_VERSION_NUMBER(major,minor,sub)`. It is useful for checking relative version numbers. For example, a user application could include the following:

```
#include "SPEX.h"
#if SPEX_VERSION < SPEX_VERSION_NUMBER(3,2,0)
// SPEX version is prior to 3.2.0
#else
// SPEX is version 3.2.0 or later
#endif
```

4.9.2 SPEX_determine_symmetry: Determine if a matrix is symmetric

```
SPEX_info SPEX_determine_symmetry
(
    bool *is_symmetric,          // true if matrix is symmetric, false otherwise
    const SPEX_matrix A,         // Input matrix to be checked for symmetry
    const SPEX_options option
) ;
```

`SPEX_determine_symmetry` checks if A has a symmetric pattern and is also numerically symmetric. If A is numerically a symmetric matrix and has a symmetric pattern, `is_symmetric` is returned as `true`. Otherwise, `is_symmetric` is returned as `false`.

4.9.3 SPEX_transpose: Transpose a CSC mpz matrix

```
SPEX_info SPEX_transpose
(
    SPEX_matrix *C_handle,       // C = A'
    SPEX_matrix A,               // Matrix to be transposed
    const SPEX_options option
) ;
```

`SPEX_transpose` sets $C = A^T$. Currently, it is only supported if A is CSC and `mpz_t`. Returns `SPEX_OK` if successful; otherwise it returns the appropriate error code.

4.10 SPEX_gmp: SPEX wrapper functions for GMP/MPFR

SPEX provides a set of memory-safe wrappers for all GMP and MPFR functions used by SPEX. The wrappers provide error-handling for out-of-memory conditions that are not handled by the GMP and MPFR libraries. These wrapper functions are used inside all SPEX functions, wherever any GMP or MPFR functions are used. These functions may also be called by the end-user application.

Each wrapped function has the same name as its corresponding GMP/MPFR function with the added prefix `SPEX_`. For example, the default GMP function `mpz_mul` is changed to `SPEX_mpz_mul`. Each SPEX GMP/MPFR function returns `SPEX_OK` if successful or the correct error code if not. The following table gives a brief list of each currently covered SPEX GMP/MPFR function. Each function is declared in `SPEX.h` and defined in `SPEX/SPEX_Util/Source/SPEX_gmp.c`.

MPFR Function	SPEX_MPFR Function	Description
<code>mpfr_asprintf(&buff, fmt, ...)</code>	<code>SPEX_mpfr_asprintf(&buff, fmt, ...)</code>	Print format to allocated string
<code>mpfr_clear(x)</code>	<code>SPEX_mpfr_clear(x)</code>	Safely free <code>mpfr_t</code> value
<code>mpfr_div_d(x, y, z, rnd)</code>	<code>SPEX_mpfr_div_d(x, y, z, rnd)</code>	$x = y/z$ (double)
<code>mpfr_free_cache()</code>	<code>SPEX_mpfr_free_cache()</code>	Free all caches and pools used by MPFR internally
<code>mpfr_free_str(buff)</code>	<code>SPEX_mpfr_free_str(buff)</code>	Free string allocated by MPFR
<code>x = mpfr_get_d(y, rnd)</code>	<code>SPEX_mpfr_get_d(x, y, rnd)</code>	(double) $x = y$
<code>mpfr_get_q(x, y)</code>	<code>SPEX_mpfr_get_q(x, y, rnd)</code>	(<code>mpq_t</code>) $x = y$
<code>x = mpfr_get_si(y, rnd)</code>	<code>SPEX_mpfr_get_si(x, y, rnd)</code>	(<code>int64_t</code>) $x = y$
<code>r = mpfr_get_z(x, y, rnd)</code>	<code>SPEX_mpfr_get_z(x, y, rnd)</code>	(<code>mpz_t</code>) $x = y$
<code>mpfr_init2(x, size)</code>	<code>SPEX_mpfr_init2(x, size)</code>	Initialize x with size bits
<code>mpfr_mul(x, y, z, rnd)</code>	<code>SPEX_mpfr_mul(x, y, z, rnd)</code>	$x = y * z$ (<code>mpfr_t</code>)
<code>mpfr_mul_d(x, y, z, rnd)</code>	<code>SPEX_mpfr_mul_d(x, y, z, rnd)</code>	$x = y * z$ (double)
<code>mpfr_set(x, y, rnd)</code>	<code>SPEX_mpfr_set(x, y, rnd)</code>	$x = y$
<code>mpfr_set_d(x, y, rnd)</code>	<code>SPEX_mpfr_set_d(x, y, rnd)</code>	$x = y$ (double)
<code>mpfr_set_null(x)</code>	<code>SPEX_mpfr_set_null(x)</code>	Initialize the (pointer) contents of a <code>mpfr_t</code> value
<code>mpfr_set_prec(x, size)</code>	<code>SPEX_mpfr_set_prec(x, size)</code>	Set the precision of an <code>mpfr_t</code> number
<code>mpfr_set_q(x, y, rnd)</code>	<code>SPEX_mpfr_set_q(x, y, rnd)</code>	$x = y$ (<code>mpq_t</code>)
<code>mpfr_set_si(x, y, rnd)</code>	<code>SPEX_mpfr_set_si(x, y, rnd)</code>	$x = y$ (<code>int64_t</code>)
<code>mpfr_set_z(x, y, rnd)</code>	<code>SPEX_mpfr_set_z(x, y, rnd)</code>	$x = y$ (<code>mpz_t</code>)
<code>sgn = mpfr_sgn(x)</code>	<code>SPEX_mpfr_sgn(sgn, x)</code>	$sgn = \text{sgn}(x)$
<code>mpfr_ui_pow_ui(x, y, z, rnd)</code>	<code>SPEX_mpfr_ui_pow_ui(x, y, z, rnd)</code>	$x = y^z$ (<code>uint64_t</code>)

GMP Function	SPEX_GMP Function	Description
<code>gmp_fscanf(fp, fmt, ...)</code>	<code>SPEX_gmp_fscanf(fp, fmt, ...)</code>	Read from file <code>fp</code>
<code>mpq_abs(x, y)</code>	<code>SPEX_mpq_abs(x, y)</code>	$x = y $
<code>mpq_add(x, y, z)</code>	<code>SPEX_mpq_add(x, y, z)</code>	$x = y + z$
<code>mpq_canonicalize(x)</code>	<code>SPEX_mpq_canonicalize(x)</code>	Canonicalize <code>x</code>
<code>mpq_clear(x)</code>	<code>SPEX_mpq_clear(x)</code>	Safely free <code>mpq_t</code> value
<code>r = mpq_cmp(x, y)</code>	<code>SPEX_mpq_cmp(r, x, y)</code>	$r = \text{sgn}(x - y)$
<code>r = mpq_cmp_ui(x, n, d)</code>	<code>SPEX_mpq_cmp_ui(r, x, n, d)</code>	$r = \text{sgn}(x - n/d)$ (<code>uint64_t/uint64_t</code>)
<code>mpq_div(x, y, z)</code>	<code>SPEX_mpq_div(x, y, z)</code>	$x = y/z$
<code>r = mpq_equal(x, y)</code>	<code>SPEX_mpq_equal(r, x, y)</code>	$r \neq 0$ if $x = y$, $r = 0$ if $x \neq y$
<code>x = mpq_get_d(y)</code>	<code>SPEX_mpq_get_d(x, y)</code>	(double) $x = y$
<code>mpq_init(x)</code>	<code>SPEX_mpq_init(x)</code>	Initialize <code>x</code>
<code>mpq_mul(x, y, z)</code>	<code>SPEX_mpq_mul(x, y, z)</code>	$x = y * z$
<code>mpq_neg(x, y)</code>	<code>SPEX_mpq_neg(x, y)</code>	$x = -y$
<code>mpq_set(x, y)</code>	<code>SPEX_mpq_set(x, y)</code>	$x = y$
<code>mpq_set_d(x, y)</code>	<code>SPEX_mpq_set_d(x, y)</code>	$x = y$ (double)
<code>mpq_set_den(x, y)</code>	<code>SPEX_mpq_set_den(x, y)</code>	$\text{den}(x) = y$
<code>mpq_set_null(x)</code>	<code>SPEX_mpq_set_null(x)</code>	Initialize the (pointer) contents of a <code>mpq_t</code> value
<code>mpq_set_num(x, y)</code>	<code>SPEX_mpq_set_num(x, y)</code>	$\text{num}(x) = y$
<code>mpq_set_si(x, y, z)</code>	<code>SPEX_mpq_set_si(x, y, z)</code>	$x = y/z$ (<code>int64_t/uint64_t</code>)
<code>mpq_set_ui(x, y, z)</code>	<code>SPEX_mpq_set_ui(x, y, z)</code>	$x = y/z$ (<code>uint64_t/uint64_t</code>)
<code>mpq_set_z(x, y)</code>	<code>SPEX_mpq_set_z(x, y)</code>	$x = y$ (<code>mpz</code>)
<code>sgn = mpq_sgn(x)</code>	<code>SPEX_mpq_sgn(sgn, x)</code>	$\text{sgn} = \text{sgn}(x)$
<code>mpz_abs(x, y)</code>	<code>SPEX_mpz_abs(x, y)</code>	$x = y $
<code>mpz_add(x, y, z)</code>	<code>SPEX_mpz_add(x, y, z)</code>	$x = y + z$
<code>mpz_addmul(x, y, z)</code>	<code>SPEX_mpz_addmul(x, y, z)</code>	$x = x + y * z$
<code>mpz_cdiv_q(q, x, y)</code>	<code>SPEX_mpz_cdiv_q(q, x, y)</code>	$q = \text{ceil}(x/y)$
<code>mpz_cdiv_qr(q, r, x, y)</code>	<code>SPEX_mpz_cdiv_qr(q, r, x, y)</code>	$q = \text{ceil}(x/y), r = x - q * y$
<code>mpz_clear(x)</code>	<code>SPEX_mpz_clear(x)</code>	Safely free <code>mpz_t</code> value
<code>r = mpz_cmp(x, y)</code>	<code>SPEX_mpz_cmp(r, x, y)</code>	$r = \text{sgn}(x - y)$
<code>r = mpz_cmp_ui(x, y)</code>	<code>SPEX_mpz_cmp_ui(r, x, y)</code>	$r = \text{sgn}(x - y)$ (<code>uint64_t</code>)
<code>r = mpz_cmpabs(x, y)</code>	<code>SPEX_mpz_cmpabs(r, x, y)</code>	$r = \text{sgn}(x - y)$
<code>r = mpz_cmpabs_ui(x, y)</code>	<code>SPEX_mpz_cmpabs_ui(r, x, y)</code>	$r = \text{sgn}(x - y)$ (<code>uint64_t</code>)
<code>mpz_divexact(x, y, z)</code>	<code>SPEX_mpz_divexact(x, y, z)</code>	$x = y/z$
<code>mpz_fdiv_q(q, x, y)</code>	<code>SPEX_mpz_fdiv_q(q, x, y)</code>	$q = \text{floor}(x/y)$
<code>gcd = mpz_gcd(x, y)</code>	<code>SPEX_mpz_gcd(gcd, x, y)</code>	$\text{gcd} = \text{gcd}(x, y)$
<code>x = mpz_get_d(y)</code>	<code>SPEX_mpz_get_d(x, y)</code>	(double) $x = y$
<code>x = mpz_get_si(y)</code>	<code>SPEX_mpz_get_si(x, y)</code>	(<code>int64_t</code>) $x = y$
<code>mpz_init(x)</code>	<code>SPEX_mpz_init(x)</code>	Initialize <code>x</code>
<code>mpz_init2(x, size)</code>	<code>SPEX_mpz_init2(x, size)</code>	Initialize <code>x</code> to <code>size</code> bits
<code>lcm = mpz_lcm(x, y)</code>	<code>SPEX_mpz_lcm(lcm, x, y)</code>	$\text{lcm} = \text{lcm}(x, y)$
<code>mpz_mul(x, y, z)</code>	<code>SPEX_mpz_mul(x, y, z)</code>	$x = y * z$
<code>mpz_mul_si(x, y, z)</code>	<code>SPEX_mpz_mul_si(x, y, z)</code>	$x = y * z$ (<code>int64_t</code>)
<code>mpz_neg(x, y)</code>	<code>SPEX_mpz_neg(x, y)</code>	$x = -y$
<code>mpz_set(x, y)</code>	<code>SPEX_mpz_set(x, y)</code>	$x = y$ (<code>mpz_t</code>)
<code>mpz_set_null(x)</code>	<code>SPEX_mpz_set_null(x)</code>	Initialize the (pointer) contents of a <code>mpz_t</code> value
<code>mpz_set_si(x, y)</code>	<code>SPEX_mpz_set_si(x, y)</code>	$x = y$ (<code>int64_t</code>)
<code>mpz_set_ui(x, y)</code>	<code>SPEX_mpz_set_ui(x, y)</code>	$x = y$ (<code>uint64_t</code>)
<code>sgn = mpz_sgn(x)</code>	<code>SPEX_mpz_sgn(sgn, x)</code>	$\text{sgn} = \text{sgn}(x)$
<code>size = mpz_sizeinbase(x, base)</code>	<code>SPEX_mpz_sizeinbase(size, x, base)</code>	size of <code>x</code> in base
<code>mpz_sub(x, y, z)</code>	<code>SPEX_mpz_sub(x, y, z)</code>	$x = y - z$
<code>mpz_submul(x, y, z)</code>	<code>SPEX_mpz_submul(x, y, z)</code>	$x = x - y * z$
<code>mpz_swap(x, y)</code>	<code>SPEX_mpz_swap(x, y)</code>	Swap the values of <code>x</code> and <code>y</code>

If additional GMP and MPFR functions are needed in the end-user application, this wrapper mechanism can be extended to those functions, which requires the source files of the SPEX library to be revised, (i.e., both `SPEX.h` and `SPEX_gmp.c`). Below are instructions on how to do this.

Consider a GMP function `void gmpfunc(TYPEa a, TYPEb b, ...)`, where `TYPEa` and `TYPEb` can be GMP type data (`mpz_t`, `mpq_t` and `mpfr_t`, for example) or non-GMP type

data (int, double, for example). A wrapper for a new GMP or MPFR function can be created by following this outline:

```

SPEX_info SPEX_gmpfunc
(
    TYPEa a,
    TYPEb b,
    ...
)
{
    // Start the GMP Wrapper
    // uncomment one of the following:
    // If this function is not modifying any GMP/MPFR type variable, use
    //SPEX_GMP_WRAPPER_START;
    // If this function is modifying mpz_t type (say TYPEa = mpz_t), use
    //SPEX_GMPZ_WRAPPER_START(a) ;
    // If this function is modifying two variables of mpz_t type (say
    // TYPEa = mpz_t, TYPEb = mpz_t), use
    //SPEX_GMPZ_WRAPPER_START2(a, b) ;
    // If this function is modifying mpq_t type (say TYPEa = mpq_t), use
    //SPEX_GMPQ_WRAPPER_START(a) ;
    // If this function is modifying mpfr_t type (say TYPEa = mpfr_t), use
    //SPEX_GMPFR_WRAPPER_START(a) ;

    // Call the GMP function
    gmpfunc(a,b,...) ;

    //Finish the wrapper and return ok if successful.
    SPEX_GMP_WRAPPER_FINISH;
    return SPEX_OK;
}

```

All of the wrapped GMP/MPFR functions always return `SPEX_info` to the caller. Therefore, for some GMP/MPFR functions that have their own return value, that value becomes a parameter instead. For example, for the GMP function `int mpq_cmp(const mpq_t a, const mpq_t b)`, the return value becomes a parameter of the wrapped function. In general, a GMP/MPFR function in the form of `TYPEr gmpfunc(TYPEa a, TYPEb b, ...)`, the wrapped function can be constructed as follows:

```

SPEX_info SPEX_gmpfunc
(
    TYPEr *r,          // return value of the GMP/MPFR function
    TYPEa a,
    TYPEb b,
    ...
)
{
    // Start the GMP Wrapper (see details above)
    //SPEX_GMP_WRAPPER_START;

    // Call the GMP function
    *r = gmpfunc(a,b,...) ;

    //Finish the wrapper and return ok if successful.
    SPEX_GMP_WRAPPER_FINISH;
    return SPEX_OK;
}

```

4.11 SPEX Helper Macros

In addition to the functionality described in this section; SPEX offers several helper macros to increase ease for the end user application. The first two macros form a simple try/catch mechanism which can be used to wrap functions for error handling. The next two give an easy interface to access individual entries in a matrix.

4.11.1 SPEX_TRY and SPEX_CATCH

In a robust application, the return values from SPEX should be checked and properly handled in the case an error occurs. SPEX is written in C and thus it cannot rely on the try/catch mechanism of C++. Thus, `SPEX_TRY` and `SPEX_CHECK` aim to provide a similar mechanism within the SPEX environment. We provide `SPEX_TRY` and leave `SPEX_CATCH` to the user to define.

```

#define SPEX_TRY(method)      \
{                               \
    SPEX_info info = (method) ; \
    if (info != SPEX_OK)       \
    {                           \
        SPEX_CATCH (info) ;    \
    }                           \
}

```

An example definition of a `SPEX_CATCH` is shown below. This example assumes that the user function needs to free a matrix and return an error code to its caller.

```

#define SPEX_CATCH(info)                                \
{                                                         \
    SPEX_matrix_free (&A, NULL) ;                       \
    fprintf (stderr, "SPEX failed: info %d, line %d, file %s\n", \
            info, __LINE__, __FILE__) ;                 \
    return (info) ;                                     \
}

```

With this mechanism, the user can safely wrap any SPEX function which returns `SPEX_info` with `SPEX_TRY`.

4.11.2 SPEX_1D: Access matrix entries with 1D linear indexing

```

#define SPEX_1D(A,k,type) ((A)->x.type [k])

```

This access the k th entry of a matrix stored in any kind (CSC, triplet, dense) of any type (`mpq_t`, `mpz_t`, `int64_t`, `double`, `int`). For example, to return the k th entry of a CSC matrix with `mpz_t` data types, one would use `SPEX_1D(A, k, mpz)`.

4.11.3 SPEX_2D: Access dense matrix entries with 2D indexing

```

#define SPEX_2D(A,i,j,type) SPEX_1D (A, (i)+(j)*((A)->m), type)

```

This accesses the (i, j) entry of a dense matrix of any type (`mpq_t`, `mpz_t`, `int64_t`, `double`, `int`). For example to return the (i, j) entry of a dense matrix with `mpq_t` data types, one would use `SPEX_2D(A, i, j, mpq)`. This macro cannot be used with sparse (CSC or dynamic CSC) or triplet matrices.

5.1 SPEX LU Module Overview

SPEX_LU is a Module designed to exactly solve unsymmetric sparse linear systems, $A\mathbf{x} = \mathbf{b}$, where $A \in \mathbb{Q}^{n \times n}$, $\mathbf{b} \in \mathbb{Q}^{n \times r}$, and $\mathbf{x} \in \mathbb{Q}^{n \times r}$, using a left-looking, roundoff-error-free (REF) LU factorization $PAQ = LDU$, where L and U are integer, D is diagonal, and P and Q are row and column permutations, respectively. Note that, in order to solve a linear system, the matrix D is never explicitly computed nor needed; thus the SPEX LU Module uses only the matrices L and U . The theory associated with this code is the Sparse Left-looking Integer-Preserving (SLIP) LU factorization [7].

Version 1.1.2 of SPEX Left LU was published in ACM TOMS as: Lourenco, C., Chen, J., Moreno-Centeno, E., & Davis, T. A. (2022). Algorithm 1021: SPEX Left LU, Exactly Solving Sparse Linear Systems via a Sparse Left-looking Integer-preserving LU Factorization. ACM Transactions on Mathematical Software (TOMS), 48(2), 1-23.

5.2 Licensing

Copyright: The copyright of this software is held by Christopher Lourenco, Jinhao Chen, Erick Moreno-Centeno, and Timothy A. Davis.

Contact Info: Contact Chris Lourenco, chrisjlourenco@gmail.com, or Tim Davis, tim-davis@aldenmath.com, davis@tamu.edu, or DrTimothyAldenDavis@gmail.com

License: The SPEX LU Module is dual licensed under the GNU General Public License version 2 or the GNU Lesser General Public License version 3. Details of this license are in SPEX/License/license.txt. For alternative licenses, contact the authors.

5.3 LU Factorization and Solve Routines

To factorize and solve a linear system $A\mathbf{x} = \mathbf{b}$ via the SPEX Left LU factorization, a user must call `analyze`, `factorize`, and `solve`. The functions are explained below.

5.3.1 SPEX_lu_analyze: symbolic analysis for LU factorization


```

SPEX_info SPEX_lu_analyze
(
    SPEX_symbolic_analysis *S_handle,    // symbolic analysis including
                                         // column permutation and nnz of L and U
    const SPEX_matrix A,                // Input matrix
    const SPEX_options option
) ;

```

`SPEX_lu_analyze` performs symbolic analysis for the SPEX LU factorization. On input, the `SPEX_symbolic_analysis` `S` that `S_handle` points to is undefined; `A` must be a square matrix of `SPEX_CSC` kind; and `option` contains any parameters (default settings are used if the `option` input is `NULL`). The `option->algo` must be either `SPEX_ALGORITHM_DEFAULT` or `SPEX_LU_LEFT`. On output, `S` contains the column preordering of `A` and estimates on the number of nonzeros in `L` and `U`. The type of column permutation can be selected by changing the value of `option->order`. COLAMD is the default fill-reducing ordering, but AMD can be effective if the nonzero pattern of `A` is mostly symmetric.

5.3.2 SPEX_lu_factorize: Compute the LU factorization of A

```

SPEX_info SPEX_lu_factorize
(
    // output:
    SPEX_factorization *F_handle,    // LU factorization
    // input:
    const SPEX_matrix A,              // matrix to be factored
    const SPEX_symbolic_analysis S,   // symbolic analysis
    const SPEX_options option
) ;

```

`SPEX_lu_factorize` performs the left-looking LU factorization. On input, the `SPEX_factorization` `F` that `F_handle` points to is undefined; `A` must be a square matrix of `SPEX_CSC` `SPEX_MPZ` format; `S` is obtained from `SPEX_lu_analyze` and contains the column ordering of `A`; and `option` contains any parameters (default settings are used if `option` is input as `NULL`). Note that `option->algo` must be either `SPEX_ALGORITHM_DEFAULT` or `SPEX_LU_LEFT`. The symbolic analysis `S` is typically constructed with the same input matrix `A`, or another matrix with the same sparsity pattern. However, the only strict requirement is that the symbolic factorization `S` is created by analyzing a matrix with the dimensions as the input matrix to the numerical factorization. If the sparsity pattern is different, the numerical factorization can still be successful but the fill-in could be excessive.

On output, `A`, `S`, and `option` are unmodified and `F` contains the SPEX LU factorization of `A`. If any error occurs, `F` is returned as `NULL`, and an appropriate error code is returned.

5.3.3 SPEX_lu_solve: solve the linear system

```

SPEX_info SPEX_lu_solve    // solves the linear system LD(-1)U x = b
(
    // Output
    SPEX_matrix *x_handle,  // rational solution to the system

```

```

// input/output:
SPEX_factorization F, // The LU factorization.
                      // Mathematically, F is unchanged. However, if F
                      // is updatable on input, it is converted to
                      // non-updatable. If F is already non-updatable,
                      // it is not modified.

// input:
const SPEX_matrix b, // right hand side vector(s)
const SPEX_options option
) ;

```

`SPEX_lu_solve` obtains the solution of `mpq_t` type to the linear system $Ax = b$ upon a successful factorization. This function may be called after a successful return from `SPEX_lu_factorize`.

On input, `SPEX_matrix x` that `x_handle` points to is undefined; `F` must be a valid LU factorization; and `b` must be a dense `mpz_t` matrix with same number of rows as `F->L`; Default settings are used if `option` is input as `NULL`. Upon successful completion, the function returns `SPEX_OK`, and `x` contains the solution of `mpq_t` type with dense format to the linear system $Ax = b$. `F` is mathematically unchanged on output. However, if `F` is updatable on input, it is converted to non-updatable. If `F` is already non-updatable, it is not modified. In case of failure, `x` is returned as `NULL` and the appropriate error code is returned.

5.3.4 `SPEX_lu_backslash`: solve a linear system using an LU factorization

```

SPEX_info SPEX_lu_backslash
(
    // Output
    SPEX_matrix *x_handle, // Final solution vector
    // Input
    SPEX_type type,        // Type of output desired. Must be
                          // SPEX_MPQ, SPEX_MPFR, or SPEX_FP64

    const SPEX_matrix A,   // Input matrix
    const SPEX_matrix b,   // Right hand side vector(s)
    const SPEX_options option
) ;

```

`SPEX_lu_backslash` solves the linear system $Ax = b$ and returns the solution as a dense matrix of `mpq_t`, `mpfr_t` or `double` entries. This function performs symbolic analysis, factorization, and solving all in one function. It can be thought of as an exact version of an LU factorization based MATLAB sparse backslash.

On input, `SPEX_matrix x` that `X_handle` points to is undefined. `type` must be one of: `SPEX_MPQ`, `SPEX_MPFR` or `SPEX_FP64` to specify the data type of the solution entries. `A` should be a square CSC `mpz_t` matrix while `b` should be a dense `mpz_t` matrix. In addition, `A->m` should be equal to `b->m`. Note that `option->algo` must be either `SPEX_ALGORITHM_DEFAULT` or `SPEX_LU_LEFT`. Default settings are used if `option` is input as `NULL`.

Upon successful completion, the function returns `SPEX_OK`, and `x` contains the solution of data type specified by `type` to the linear system $Ax = b$. In case of failure, `x` is returned as `NULL` and the appropriate error code is returned.

CHAPTER 6

SPEX CHOLESKY AND LDL MODULE

6.1 SPEX Cholesky and LDL Module Overview

The SPEX Cholesky Module is designed to exactly solve symmetric linear systems, $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{Q}^{n \times n}$, $\mathbf{b} \in \mathbb{Q}^{n \times r}$, and $\mathbf{x} \in \mathbb{Q}^{n \times r}$, using either an exact Cholesky factorization (appropriate for symmetric positive definite matrices) and an exact LDL factorization (appropriate for symmetric matrices with non-singular leading principle minors). This package performs either a left-looking or up-looking sparse roundoff-error-free Cholesky or LDL factorization $PAP^T = LDL^T$ where L is integer, and P is the symmetric permutation.

Note that, in order to solve a linear system, the matrix D is never explicitly computed nor needed; thus this package uses only the matrix L . The theory associated with this code can be found at [8].

6.2 Licensing

Copyright: The copyright of this software is held by Christopher Lourenco, Lorena Mejia Domenzain, Jinhao Chen, Erick Moreno-Centeno, and Timothy A. Davis.

Contact Info: Contact Chris Lourenco, chrisjlourenco@gmail.com, or Tim Davis, tim-davis@aldenmath.com, davis@tamu.edu, or DrTimothyAldenDavis@gmail.com

License: This software package is dual licensed under the GNU General Public License version 2 or the GNU Lesser General Public License version 3. Details of this license are in `SPEX/License/license.txt`. For alternative licenses, contact the authors.

6.3 Cholesky/LDL Factorization and Solve Routines

To factorize and solve a linear system $A\mathbf{x} = \mathbf{b}$ via the SPEX Cholesky or SPEX LDL factorization, a user must call the corresponding analyze, factorize, and solve methods. There are four user callable functions for each SPEX Cholesky and SPEX LDL factorization. The functions are explained below.

6.3.1 SPEX_cholesky_analyze: symbolic analysis for Cholesky factorization

```

SPEX_info SPEX_cholesky_analyze
(
    // Output
    SPEX_symbolic_analysis *S_handle, // Symbolic analysis data structure
    // Input
    const SPEX_matrix A,             // Input matrix. Must be SPEX_MPZ and SPEX_CSC
    const SPEX_options option
) ;

```

`SPEX_cholesky_analyze` performs symbolic analysis for the SPEX Cholesky factorization. On input, the `SPEX_symbolic_analysis` `S` that `S_handle` points to is undefined; `A` must be a symmetric positive definite matrix of `SPEX_CSC` kind; and `option` contains any parameters (default settings are used if `option` is input as `NULL`). `option->algo` must be either `SPEX_ALGORITHM_DEFAULT`, `SPEX_CHOL_LEFT`, or `SPEX_CHOL_UP`. On output, `S` contains the row and column ordering of `A`, the exact number of nonzeros in `L`, the elimination tree of `A`, and the column pointers of `L`. The type of ordering can be chosen with `option->order`. AMD is the default ordering, but sometimes COLAMD can give good results.

6.3.2 SPEX_cholesky_factorize: Compute the Cholesky factorization of A

```

SPEX_info SPEX_cholesky_factorize
(
    // Output
    SPEX_factorization *F_handle, // Cholesky factorization struct
    //Input
    const SPEX_matrix A,          // Matrix to be factored. Must be SPEX_MPZ
                                // and SPEX_CSC
    const SPEX_symbolic_analysis S, // Symbolic analysis struct containing the
                                // elimination tree of A, the column
                                // pointers of L, and the exact number of
                                // nonzeros of L.
    const SPEX_options option
) ;

```

`SPEX_cholesky_factorize` performs the SPEX Cholesky factorization via either the up-looking (default) or left-looking algorithm (specified by `option->algo`). On input, the `SPEX_factorization` `F` that `F_handle` points to is undefined; `A` must be a symmetric positive definite matrix of `SPEX_CSC` `SPEX_MPZ` format; `S` is obtained from `SPEX_cholesky_analyze` that contains the column/row ordering of `A`; and `option` contains any parameters (default settings are used if `option` is input as `NULL`). On output, `A`, `S`, and `option` are unmodified and `F` contains the SPEX Cholesky factorization of `A`.

The symbolic analysis `S` must be obtained by a call to `SPEX_cholesky_analyze` with an input matrix `A` that has the identical sparsity pattern as the matrix given as input to this `SPEX_cholesky_factorize` method. Results are undefined if this condition does not hold.

If error occurs, `F` is returned as `NULL`, and an appropriate error code is returned.

6.3.3 SPEX_cholesky_solve: solve the linear system

```

SPEX_info SPEX_cholesky_solve
(
    // Output
    SPEX_matrix *x_handle,      // On input: undefined.
                                // On output: Rational solution (SPEX_MPQ)
                                // to the system.

    // input/output:
    SPEX_factorization F,      // The Cholesky factorization.
                                // Mathematically, F is unchanged. However, if F
                                // is updatable on input, it is converted to
                                // non-updatable. If F is already non-updatable,
                                // it is not modified.

    // input:
    const SPEX_matrix b,        // Right hand side vector
    const SPEX_options option
) ;

```

`SPEX_cholesky_solve` obtains the solution of `mpq_t` type to the linear system $A\mathbf{x} = \mathbf{b}$. This function may be called after a successful return from `SPEX_cholesky_factorize`.

On input, `SPEX_matrix` `x` that `x_handle` points to is undefined; `F` must be a valid Cholesky factorization and `b` must be dense `mpz_t` matrix with the same number of rows as `F`→`L`. Default settings are used if `option` is input as `NULL`. Upon successful completion, the function returns `SPEX_OK`, and `x` contains the solution of `mpq_t` type with dense format to the linear system $A\mathbf{x} = \mathbf{b}$. `F` is mathematically unchanged on output. However, if `F` is updatable on input, it is converted to non-updatable. If `F` is already non-updatable, it is not modified. In case of failure, `x` is returned as `NULL` and the appropriate error code is returned.

6.3.4 SPEX_cholesky_backslash: solve a linear system using a Cholesky factorization

```

SPEX_info SPEX_cholesky_backslash
(
    // Output
    SPEX_matrix *x_handle,      // On input: undefined.
                                // On output: solution vector(s)

    // Input
    SPEX_type type,             // Type of output desired
                                // Must be SPEX_FP64, SPEX_MPFR, or SPEX_MPQ

    const SPEX_matrix A,        // Input matrix. Must be SPEX_MPZ and SPEX_CSC
    const SPEX_matrix b,        // Right hand side vector(s). Must be
                                // SPEX_MPZ and SPEX_DENSE

    const SPEX_options option
) ;

```

`SPEX_cholesky_backslash` solves the linear system $A\mathbf{x} = \mathbf{b}$ and returns the solution as a dense matrix of `mpq_t`, `mpfr_t` or `double` entries. This function performs symbolic analysis, factorization, and solving all in one function. It can be thought of as an exact version of the MATLAB sparse backslash for symmetric positive-definite matrices. If `A` is not symmetric positive-definite, this function will return an appropriate error code.

On input, `SPEX_matrix x` that `x_handle` points to is undefined. `type` must be one of: `SPEX_MPQ`, `SPEX_MPFR` or `SPEX_FP64` to specify the data type of the solution entries. `A` should be a square CSC `mpz_t` matrix while `b` should be a dense `mpz_t` matrix. In addition, `A->m` should be equal to `b->m`. Note that `option->algo` must be either `SPEX_ALGORITHM_DEFAULT`, `SPEX_CHOL_LEFT`, or `SPEX_CHOL_UP`. Default settings are used if `option` is input as `NULL`.

Upon successful completion, the function returns `SPEX_OK`, and `x` contains the solution of data type specified by `type` to the linear system $Ax = b$. In case of failure, `x` is returned as `NULL` and the appropriate error code is returned.

6.3.5 SPEX_ldl_analyze: symbolic analysis for LDL factorization

```
SPEX_info SPEX_ldl_analyze
(
    // Output
    SPEX_symbolic_analysis *S_handle, // Symbolic analysis data structure
    // Input
    const SPEX_matrix A,             // Input matrix. Must be SPEX_MPZ and SPEX_CSC
    const SPEX_options option
);
```

`SPEX_ldl_analyze` performs symbolic analysis for the SPEX LDL factorization. On input, the `SPEX_symbolic_analysis S` that `S_handle` points to is undefined; `A` must be a symmetric matrix of `SPEX_CSC` kind with a zero-free diagonal; and `option` contains the parameters (default settings are used if `option` is input as `NULL`). `option->algo` must be either `SPEX_ALGORITHM_DEFAULT`, `SPEX_LD_L`, or `SPEX_LD_L`. On output, `S` contains the row and column ordering of `A`, the exact number of nonzeros in `L`, the elimination tree of `A`, and the column pointers of `L`. The type of ordering can be chosen with `option->order`. AMD is the default ordering, but sometimes COLAMD can give good results.

6.3.6 SPEX_ldl_factorize: Compute the LDL factorization of A

```
SPEX_info SPEX_ldl_factorize
(
    // Output
    SPEX_factorization *F_handle, // Cholesky factorization struct
    //Input
    const SPEX_matrix A,          // Matrix to be factored. Must be SPEX_MPZ
                                // and SPEX_CSC
    const SPEX_symbolic_analysis S, // Symbolic analysis struct containing the
                                // elimination tree of A, the column
                                // pointers of L, and the exact number of
                                // nonzeros of L.
    const SPEX_options option
);
```

`SPEX_ldl_factorize` performs the SPEX LDL factorization via either the up-looking or left-looking algorithm (based on the choice of `option->algo`). On input, the `SPEX_factorization F` that `F_handle` points to is undefined; `A` must be a symmetric matrix of `SPEX_CSC` `SPEX_MPZ`

format with a zero-free diagonal; **S** is obtained from `SPEX_ldl_analyze` that contains the column/row ordering of **A**; and **option** contains any parameters (default settings are used if **option** is input as `NULL`). On output, **A**, **S**, and **option** are unmodified and **F** contains the SPEX LDL factorization of **A**.

The symbolic analysis **S** must be obtained by a call to `SPEX_ldl_analyze` with an input matrix **A** that has the identical sparsity pattern as the matrix given as input to this `SPEX_ldl_factorize` method. Results are undefined if this condition does not hold.

If error occurs, **F** is returned as `NULL`, and an appropriate error code is returned.

6.3.7 SPEX_ldl_solve: solve the linear system

```
SPEX_info SPEX_ldl_solve
(
    // Output
    SPEX_matrix *x_handle,      // On input: undefined.
                                // On output: Rational solution (SPEX_MPQ)
                                // to the system.

    // input/output:
    SPEX_factorization F,      // The LDL factorization.
                                // Mathematically, F is unchanged. However, if F
                                // is updatable on input, it is converted to
                                // non-updatable. If F is already non-updatable,
                                // it is not modified.

    // input:
    const SPEX_matrix b,        // Right hand side vector
    const SPEX_options option
);
```

`SPEX_ldl_solve` obtains the solution of `mpq_t` type to the linear system $A\mathbf{x} = \mathbf{b}$. This function may be called after a successful return from `SPEX_ldl_factorize`.

On input, `SPEX_matrix` **x** that **x_handle** points to is undefined; **F** must be a valid LDL factorization and **b** must be dense `mpz_t` with same number of rows as **F**->**L**. Default settings are used if **option** is input as `NULL`. Upon successful completion, the function returns `SPEX_OK`, and **x** contains the solution of `mpq_t` type with dense format to the linear system $A\mathbf{x} = \mathbf{b}$. **F** is mathematically unchanged on output. However, if **F** is updatable on input, it is converted to non-updatable. If **F** is already non-updatable, it is not modified. In case of failure, **x** is returned as `NULL` and the appropriate error code is returned.

6.3.8 SPEX_ldl_backslash: solve a linear system using an LDL factorization

```

SPEX_info SPEX_ldl_backslash
(
    // Output
    SPEX_matrix *x_handle,      // On input: undefined.
                                // On output: solution vector(s)

    // Input
    SPEX_type type,             // Type of output desired
                                // Must be SPEX_FP64, SPEX_MPFR, or SPEX_MPQ

    const SPEX_matrix A,        // Input matrix. Must be SPEX_MPZ and SPEX_CSC
    const SPEX_matrix b,        // Right hand side vector(s). Must be
                                // SPEX_MPZ and SPEX_DENSE

    const SPEX_options option
) ;

```

`SPEX_ldl_backslash` solves the linear system $A\mathbf{x} = \mathbf{b}$ and returns the solution as a dense matrix of `mpq_t`, `mpfr_t` or `double` entries. This function performs symbolic analysis, factorization, and solving all in one call. If A is not symmetric with non-singular leading principle minors, this function will return an appropriate error code. In this case, an LU factorization should be used instead.

On input, `SPEX_matrix x` that `x_handle` points to is undefined. `type` must be one of: `SPEX_MPQ`, `SPEX_MPFR` or `SPEX_FP64` to specify the data type of the solution entries. A should be a square CSC `mpz_t` matrix while $\mathbf{b}$ should be a dense `mpz_t` matrix. In addition, `A->m` should be equal to `b->m`. Note that `option->algo` must be either `SPEX_ALGORITHM_DEFAULT`, `SPEX_LDL_LEFT`, or `SPEX_LDL_UP`. Default settings are used if `option` is input as `NULL`.

Upon successful completion, the function returns `SPEX_OK`, and $\mathbf{x}$ contains the solution of data type specified by `type` to the linear system $A\mathbf{x} = \mathbf{b}$. In case of failure, $\mathbf{x}$ is returned as `NULL` and the appropriate error code is returned.

CHAPTER 7

SPEX BACKSLASH MODULE

7.1 SPEX Backslash Module Overview

SPEX Backslash is a SPEX Module designed to exactly solve sparse linear systems, $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{Q}^{n \times n}$, $\mathbf{b} \in \mathbb{Q}^{n \times r}$, and $\mathbf{x} \in \mathbb{Q}^{n \times r}$. This package acts as either a wrapper to the SPEX LU, Cholesky, and LDL backslash methods, or it can determine the appropriate factorization to apply based on the structure of the input matrix.

7.2 Licensing

Copyright: The copyright of this software is held by Christopher Lourenco, Lorena Mejia Domenzain, Jinhao Chen, Erick Moreno-Centeno, and Timothy A. Davis.

Contact Info: Contact Chris Lourenco, chrisjlourenco@gmail.com, or Tim Davis, tim-davis@aldenmath.com, davis@tamu.edu, or DrTimothyAldenDavis@gmail.com

License: This software package is dual licensed under the GNU General Public License version 2 or the GNU Lesser General Public License version 3. Details of this license are in SPEX/License/license.txt. For alternative licenses, contact the authors.

7.3 SPEX_backslash: exactly solve a sparse linear system

```
SPEX_info SPEX_backslash
(
    // Output
    SPEX_matrix *x_handle,      // On output: Final solution vector(s)
                                // On input: undefined

    // Input
    const SPEX_type type,      // Type of output desired
                                // Must be SPEX_MPQ, SPEX_MPFR, or SPEX_FP64

    const SPEX_matrix A,       // Input matrix
    const SPEX_matrix b,       // Right hand side vector(s)
    const SPEX_options option

) ;
```

`SPEX_backslash` exactly solves the linear system $A\mathbf{x} = \mathbf{b}$ using the appropriate factorization. On input, `SPEX_matrix x` that `X_handle` points to is undefined. `type` must be one of: `SPEX_MPQ`, `SPEX_MPFR` or `SPEX_FP64` to specify the data type of the solution entries. `A` should be a square CSC `mpz_t` matrix while `b` should be a dense `mpz_t` matrix. In addition, `A->m` should be equal to `b->m`. Default settings are used if `option` is input as `NULL`.

The behavior of `SPEX_Backslash` depends on the `option->algo` parameter. If `option->algo` is left as default: the symmetry of A is checked. If A is symmetric, an up-looking LDL factorization is attempted. If the factorization succeeds, `x` is returned. Otherwise, a left-looking LU factorization is attempted. If `option->algo` is not set to `SPEX_ALGORITHM_DEFAULT` the request algorithm requested is used with no substitutions. For example, if `option->algo` is set to `SPEX_CHOL_LEFT`, then `SPEX_Backslash` essentially serves as a direct wrapper to the SPEX left-looking Cholesky factorization method, `SPEX_cholesky_backslash`. An appropriate error code is returned if the selected algorithm is not appropriate for the given matrix (e.g., if `option->algo` is `SPEX_LDL_LEFT` but the matrix is not symmetric).

Upon successful completion, the function returns `SPEX_OK`, and `x` contains the solution of data type specified by `type` to the linear system $A\mathbf{x} = \mathbf{b}$. If an error occurs, `x` is returned as `NULL` in `X_handle`, and the appropriate error code is returned.

CHAPTER 8

SPEX UPDATE MODULE

FIXME: split this into SPEX LU ColRep and Chol rank1 Modules.

8.1 SPEX Update Module Overview

SPEX Update is a SPEX Module designed to exactly update an available SPEX LU, Cholesky, or LDL factorization for the sparse matrix A when a low-rank modification is performed on A . The low-rank modification can be either column/row replacement for an LU factorization, or a rank-1 update/downdate for a Cholesky or LDL factorization. The theoretical basis for this code is to be submitted as the following papers.

- FIXME, cite TOMS here: J. Chen, T. A. Davis, “Sparse Exact LU Update for Column Replacement,” in preparation.
- FIXME, cite TOMS here: J. Chen, T. A. Davis, “Sparse Exact Rank-1 Cholesky Update,” in preparation.

8.2 Licensing

Copyright: The copyright of this software is held by Jinhao Chen, Christopher Lourenco, Lorena Mejia Domenzain, Erick Moreno-Centeno, and Timothy A. Davis.

Contact Info: Contact Jinhao Chen cjh10644@hotmail.com, or Chris Lourenco, chrisjlourenco@gmail.com or Tim Davis, timdavis@aldenmath.com, davis@tamu.edu, or DrTimothyAldenDavis@gmail.com

License: This software package is dual licensed under the GNU General Public License version 2 or the GNU Lesser General Public License version 3. Details of this license are in `SPEX/License/license.txt`. For alternative licenses, contact the authors.

8.3 SPEX Factorization Update Functions

SPEX update provides the following functions to perform factorization updates.

8.3.1 SPEX_update_cholesky_rank1: rank-1 Cholesky update/down-date

```

SPEX_info SPEX_update_cholesky_rank1
(
    SPEX_factorization F,    // The SPEX Cholesky factorization of A, including
                            // L, rhos, P and Pinv. This factorization will be
                            // modified during the update process. Therefore,
                            // if this function fails for any reason, the
                            // returned F should be considered as undefined.
    SPEX_matrix w,          // a n-by-1 dynamic_CSC matrix that contains the
                            // vector to modify the original matrix A, the
                            // resulting A is A+sigma*w*w^T. A->scale = w->scale
                            // and w->v[0]->scale = 1. In output, w is
                            // updated as the solution to L*D^(-1)*w_out = w
    const int64_t sigma,    // a nonzero scalar that determines whether
                            // this is an update (sigma > 0) or downdate
                            // (sigma < 0).
    const SPEX_options option
) ;

```

`SPEX_update_cholesky_rank1` performs the rank-1 Cholesky update/downdate for a n -by- n sparse matrix A . On input, F contains the valid Cholesky factorization to be updated that consists of L , rhos , P_{perm} and $P_{\text{inv_perm}}$; w must be a n -by-1 matrix of `SPEX_DYNAMIC_CSC` kind; sigma is a nonzero scalar, where $\text{sigma} > 0$ indicates update and $\text{sigma} < 0$ indicates downdate; and option contains parameters (default settings are used if option is input as `NULL`).

This function requires that the rows of w are in the same order as that of A (and $F \rightarrow L$ if F is updatable). In addition, w and `SPEX_matrix` A that contains A are scaled with same factor, i.e., $w \rightarrow \text{scale} = A \rightarrow \text{scale}$, while $w \rightarrow v[0] \rightarrow \text{scale} = 1$. (Users can refer to section 3.5 for the difference between $w \rightarrow \text{scale}$ and $w \rightarrow v[0] \rightarrow \text{scale}$.) As for F , if it is input in a non-updatable format, it is first converted to updatable format.

On success `SPEX_OK` is returned, and F contains the updatable roundoff-error-free Cholesky factorization of $A + \sigma w w^T$, and w contains the solution w_{out} to $LD^{-1}w_{\text{out}} = w_{\text{in}}$, where $A = LD^{-1}L^T$ and w_{in} is the input matrix w . It should be noted that w_{out} is also the solution to $\bar{L}\bar{D}^{-1}w_{\text{out}} = w_{\text{in}}$, where $A + \sigma w w^T = \bar{L}\bar{D}^{-1}\bar{L}^T$. Otherwise, if this function fails for any reason, the appropriate error code is returned and F is undefined on output (since the modification for the factorization is done in place during the update process).

This function does not require/use/modify the original matrix A . Therefore, if the updated matrix $\bar{A}$ is needed, users should implement it by their own **BEFORE** using this function (since w will be modified by this function).

8.3.2 SPEX_update_lu_colrep: Column-replacement LU update

```

SPEX_info SPEX_update_lu_colrep
(
    SPEX_factorization F,    // The SPEX factorization of A, including L, U,
                            // rhos, P, Pinv, Q and Qinv. The factorization
                            // will be modified during the update process.
                            // Therefore, if this function fails for any
                            // reason, the returned F should be considered as
                            // undefined.

    // todo: decide on utilities to help create this n-by-1 matrix:
    SPEX_matrix vk,          // Pointer to a n-by-1 dynamic_CSC matrix
                            // which contains the column to be inserted.
                            // vk->scale = A->scale and vk->v[0]->scale = 1.
                            // The rows of vk are in the same order as A.

    int64_t k,              // The column index that vk will be inserted, 0<=k<n

    const SPEX_options option
) ;

```

`SPEX_update_lu_colrep` exactly updates the roundoff-error-free LU factorization of a n -by- n sparse matrix A when one column of A is replaced. On input F contains a valid SPEX LU factorization to be updated that consists of L , U , rhos , P_{perm} , $P_{\text{inv_perm}}$ and Q_{perm} (and $Q_{\text{inv_perm}}$ for an updatable F); vk is a n -by-1 `SPEX_DYNAMIC_CSC` matrix that contains the column to be inserted; k is the index of the column of A to be replaced and thus must be in the range of $[0, n - 1]$; and `option` contains parameters (default settings are used if `option` is input as `NULL`).

This function requires that the rows of vk are in the same permutation order as that of A (and $F \rightarrow L$ if F is updatable). In addition, vk and `SPEX_matrix` A that contains A are scaled with same factor, i.e., $vk \rightarrow \text{scale} = A \rightarrow \text{scale}$, while $vk \rightarrow v[0] \rightarrow \text{scale} = 1$. Refer to section 3.5 for the difference between $vk \rightarrow \text{scale}$ and $vk \rightarrow v[0] \rightarrow \text{scale}$. As for F , if it is input in a non-updatable format, it is first converted to updatable format.

On success `SPEX_OK` is returned, and F contains the updatable roundoff-error-free LU factorization of $\bar{A}$ which is obtained by replacing column k of A with vk , while vk is unmodified. Otherwise, if this function fails for any reason, the appropriate error code is returned and F is undefined on output (since the modification for the factorization is done in place during the update process).

This function does not require/use/modify the original matrix A . Therefore, if the updated matrix $\bar{A}$ is needed, users should call `SPEX_update_matrix_colrep` **AFTER** using this function (since vk is be modified by `SPEX_update_matrix_colrep`).

FIXME: I'm confused; the above paragraph says the functions should be called in the order: `SPEX_update_lu_colrep` and then `SPEX_update_matrix_colrep`. But this is inconsistent with the discussion in Section 8.5.1. Which is it?

8.4 SPEX Update Solve Routines

When factorization is updated iteratively, it should not be converted back and forth between updatable and non-updatable formats. However, this happens when users use the default LU/Cholesky solves (i.e., `SPEX_lu_solve` or `SPEX_cholesky_solve`) after using either

SPEX_update_cholesky_rank1 or SPEX_update_lu_colrep, since the default LU/Cholesky solves always convert factorization to non-updatable format while the SPEX Update functions require and output updatable factorization. To avoid this, the SPEX Update Module provides functions to solve the linear equation $Ax = b$ based on the updatable LU/Cholesky factorization of A . In particular, both a normal solve and a transpose solve using the updatable factorizations are provided.

8.4.1 SPEX_update_solve: Solve linear system with the updatable factorization

```
SPEX_info SPEX_update_solve // solves  $Ax = b$  via LU or Cholesky factorization
(
    // Output
    SPEX_matrix *x_handle, // a m*n dense matrix contains the solution to
                          // the system.

    // input:
    SPEX_factorization F, // The SPEX LU or Cholesky factorization
    const SPEX_matrix b,  // a m*n dense matrix contains the right-hand-side
                          // vector
    const SPEX_options option
) ;
```

SPEX_update_solve obtains the solution of `mpq_t` type to the linear system $Ax = b$ upon a valid LU/Cholesky factorization of A .

On input, `SPEX_matrix x` that `x_handle` points to is undefined; F must be a valid LU/Cholesky factorization; b must be dense `mpz_t` with same number of rows as $F \rightarrow L$; and default settings are used if `option` is input as `NULL`. Upon successful completion, the function returns `SPEX_OK`, and the dense `mpq_t` matrix x contains the solution to the linear system $Ax = b$. F is mathematically unchanged on output. However, if F is non-updatable on input, it is converted to updatable. If F is already updatable, it is not modified. In case of failure, x is returned as `NULL` and the appropriate error code is returned.

8.4.2 SPEX_update_tsolve: Solve transposed linear system with the updatable factorization

```
SPEX_info SPEX_update_tsolve // solves  $A^T x = b$ 
(
    // Output
    SPEX_matrix *x_handle, // a m*n dense matrix contains the solution to
                          // the system.

    // input:
    SPEX_factorization F, // The SPEX LU or Cholesky factorization of A
    const SPEX_matrix b,  // a m*n dense matrix contains the right-hand-side
                          // vector
    const SPEX_options option
) ;
```

SPEX_update_tsolve obtains the solution of `mpq_t` type to the linear system $A^T x = b$ upon a valid LU/Cholesky factorization of A .

On input, `SPEX_matrix x` that `x_handle` points to is undefined; `F` must be a valid LU/C-holesky factorization; `b` must be dense `mpz_t` with same number of rows as `F->L`; and default settings are used if `option` is input as `NULL`. Upon successful completion, the function returns `SPEX_OK`, and the dense `mpz_t` matrix `x` contains the solution to the linear system $A^T x = b$. `F` is mathematically unchanged on output. However, if `F` is non-updatable on input, it is converted to updatable. If `F` is already updatable, it is not modified. In case of failure, `x` is returned as `NULL` and the appropriate error code is returned.

8.5 SPEX Update Helper Function

8.5.1 `SPEX_update_matrix_colrep`: Swap the columns of a matrix

```
SPEX_info SPEX_update_matrix_colrep // performs column replacement
(
    SPEX_matrix A,           // m-by-n target matrix of SPEX_DYNAMIC_CSC MPZ
    SPEX_matrix vk,          // m-by-1 SPEX_DYNAMIC_CSC MPZ matrix that contains
                            // the column vector to replace the k-th column of A
                            // vk->scale = A->scale and vk->v[0]->scale = 1.
    int64_t k,               // The column index that vk will be inserted, 0<=k<n
    const SPEX_options option
) ;
```

`SPEX_update_matrix_colrep` serves as a helper function in the column-replacement scenario where users wish to obtain the matrix after one of its columns is replaced. It should be noted that this does **NOT** update the factorization of `A`, and thus users should call `SPEX_update_lu_colrep` to update the SPEX LU factorization for column replacement. To be more specific, this function swaps column `k` of the matrix `A` with column of `vk` (i.e., swaps `A->v[k]` and `vk->v[0]`). Therefore, both `A` and `vk` are modified by this function, which requires users to call this function **BEFORE** using `SPEX_update_lu_colrep`.

FIXME: I'm confused; the above paragraph says the functions should be called in the order: `SPEX_update_matrix_colrep` and then `SPEX_update_lu_colrep`. But this is inconsistent with the discussion in Section 8.3.2. Which is it?

On input, `A` and `vk` must be matrices of `SPEX_DYNAMIC_CSC` kind with the same number of rows `m`; `k` is the index of the column of `A` to be replaced and thus must be in the range of $[0, m - 1]$; and `option` contains parameters (default settings are used if `option` is input as `NULL`).

This function requires that the rows of `vk` are in the same order as that of `A`. In addition, `vk` and `A` are scaled with the same factor, i.e., `vk->scale = A->scale`, while `vk->v[0]->scale = 1`. (Users can refer to section 3.5 for the difference between `vk->scale` and `vk->v[0]->scale`.) Different from `SPEX_update_lu_colrep`, this function does not require `A` to be a square matrix, as long as it has the same number of rows as `vk`.

On output, `A->v[k]` and `vk->v[0]` are simply swapped.

FIXME: Tim is HERE:

The MATLAB interface of SPEX can be installed by navigating to the `MATLAB` folder and typing `spex_mex_install` in the MATLAB Command Window. Doing so installs SPEX and allows the use of four MATLAB functions: `spex_lu_backslash.m`, `spex_cholesky_backslash.m`, `spex_ldl_backslash.m` and `spex_backslash.m`. Section 9.1 describes the `option` input parameter for each method. The use of each SPEX method is discussed in Section 9.2. The `SPEX/SPEX/MATLAB` folder must be in your MATLAB path.

9.1 Optional parameter settings

The SPEX MATLAB interface includes an `option` struct as an optional input parameter that modifies the behavior of each method. If this parameter is not provided, default parameter settings are used. The elements of the `option` struct are listed below. Any fields not present in the struct are treated as their default values.

- `option.pivot`: This parameter is a string that controls the pivoting scheme used. When selecting a pivot entry in a given column, the factorization method uses one of the following pivoting strategies. This parameter only applies to LU factorization:
 - `'smallest'`: (default) smallest pivot,
 - `'diagonal'`: diagonal pivot if possible, otherwise smallest pivot,
 - `'first'`: first nonzero pivot in each column,
 - `'tol smallest'`: diagonal pivot with a tolerance (`option.tol`) for the smallest pivot,
 - `'tol largest'`: diagonal pivot with a tolerance (`option.tol`) for the largest pivot,
 - `'largest'`: largest pivot.
- `option.order`: This parameter is a string controls the fill-reducing column preordering used.
 - `'none'`: no column ordering; factorize **A** as-is.

- `'colamd'`: COLAMD ordering (default for LU)
- `'amd'`: AMD ordering (default for Cholesky and LDL)
- `option.tol`: This parameter determines the tolerance used if one of the threshold pivoting schemes is chosen. The default value is 1 and this parameter can take any value in the range $(0, 1]$. This is only valid for LU factorization.
- `option.solution`: a string determining how `x` is to be returned:
 - `'double'`: `x` is converted to a 64-bit floating-point approximate solution. This is the default.
 - `'vpa'`: `x` is returned as a `vpa` array with `option.digits` digits (default is given by the MATLAB `digits` function). The result may be inexact, if an entry in `x` cannot be represented in the specified number of digits. To convert this `x` to double, use `x=double(x)`.
 - `'char'`: `x` is returned as a cell array of strings, where `x {i} = 'numerator/denominator'` and both `numerator` and `denominator` are arbitrary-length strings of decimal digits. The result is always exact, although `x` cannot be directly used in MATLAB for numerical calculations. It can be inspected or analyzed using MATLAB string manipulation. To convert `x` to `vpa`, use `x=vpa(x)`. To convert `x` to double, use `x=double(vpa(x))`.
- `option.digits`: the number of decimal digits to use for `x`, if `option.solution` is `'vpa'`. Must be in range 2 to 2^{29} .
- `option.print`: display the inputs and outputs (0: nothing (default), 1: just errors, 2: terse, 3: all).

9.2 SPEX MATLAB functions

9.2.1 `spex_lu_backslash.m`: solve a linear system via exact LU factorization

The `spex_lu_backslash.m` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times m}$ and $\mathbf{b} \in \mathbb{R}^{n \times m}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision.

This function expects as input a matlab matrix `A` and dense set of right hand side vectors `b`. `A` does not necessarily have to be sparse even though SPEX will internally treat it as so. Optionally, `option` struct can be passed in. Currently, there are 2 ways to use this function outlined below.

- `x = spex_lu_backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution vectors are more accurate than the solution obtained via `x = A \setminus b`. The solution `x` is returned as a MATLAB double matrix.

- `x = spex_lu_backslash(A,b,option)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` struct.

If the result `x` is held as a MATLAB double matrix, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution `x` may also be returned as a MATLAB `vpa` array, or as a cell array of strings; See Section 9.1 for details.

9.2.2 `spex_cholesky_backslash.m`: solve a linear system via exact Cholesky factorization

The `spex_cholesky_backslash.m` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times m}$ and $\mathbf{b} \in \mathbb{R}^{n \times m}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision. Note that A must be SPD otherwise this function returns an error.

This function expects as input a matrix A and dense set of right hand side vectors b . A does not necessarily have to be sparse even though SPEX will internally treat it as so. Optionally, the `option` struct can be passed in. Currently, there are 2 ways to use this function outlined below.

- `x = spex_cholesky_backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution vectors are more accurate than the solution obtained via `x = A \setminus b`. The solution `x` is returned as a MATLAB double matrix.
- `x = spex_cholesky_backslash(A,b,option)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` struct.

If the result `x` is held as a MATLAB double matrix, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution `x` may also be returned as a MATLAB `vpa` array, or as a cell array of strings; See Section 9.1 for details.

9.2.3 `spex_ldl_backslash.m`: solve a linear system via exact LDL factorization

The `spex_ldl_backslash.m` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times m}$ and $\mathbf{b} \in \mathbb{R}^{n \times m}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision. Note that all leading principal minors of A must be nonzero. Otherwise, this function returns an error.

This function expects as input a matrix A and dense set of right hand side vectors b . A does not necessarily have to be sparse even though SPEX will internally treat it as so. Optionally, the `option` struct can be passed in. Currently, there are 2 ways to use this function outlined below.

- `x = spex_ldl_backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution vectors are more accurate than the solution obtained via `x = A \ b`. The solution `x` is returned as a MATLAB double matrix.
- `x = spex_ldl_backslash(A,b,option)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` struct.

If the result `x` is held as a MATLAB double matrix, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution `x` may also be returned as a MATLAB `vpa` array, or as a cell array of strings; See Section 9.1 for details.

9.2.4 `spex_backslash.m`: solve a linear system via an exact factorization

The `spex_backslash.m` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times m}$ and $\mathbf{b} \in \mathbb{R}^{n \times m}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision.

This function expects as input a matrix A and dense set of right hand side vectors $\mathbf{b}$. A does not necessarily have to be sparse even though SPEX will internally treat it as so. Optionally, the `option` struct can be passed in. If A is numerically symmetric, it attempts an LDL factorization. If the LDL fails or if the matrix is not numerically symmetric it performs an LU factorization. Currently, there are 2 ways to use this function outlined below.

- `x = spex_backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution vectors are more accurate than the solution obtained via `x = A \ b`. The solution `x` is returned as a MATLAB double matrix.
- `x = spex_backslash(A,b,option)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` struct.

If the result `x` is held as a MATLAB double matrix, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution `x` may also be returned as a MATLAB `vpa` array, or as a cell array of strings; See Section 9.1 for details.

9.2.5 `spex_mex_demo.m`: example of how to use SPEX in MATLAB

This function provides a demo of the SPEX library. It shows the usage for an exact solution as well as error checking and tuning the parameters. The typical output of this function may be seen in the provided `MATLAB/html` folder.

CHAPTER 10

USING SPEX IN PYTHON

The Python interface of SPEX can be installed by navigating to the Python folder and typing `make`. Doing so allows the use of the Python SPEX library. First, this section describes the `Option` object in Section 10.1. The use of SPEX to solve $A\mathbf{x} = \mathbf{b}$ is discussed in Section 10.2.

10.1 Optional parameter settings

The SPEX Python interface includes an object as an optional input parameter that modifies behaviour. If this is not provided, default parameter settings are used.

- `output`: This parameter is a string that determines how the solution is to be returned
 - `'double'`: $\mathbf{x}$ is converted to a 64-bit floating-point approximate solution. This is the default.
 - `'string'`: $\mathbf{x}$ is returned as an array of strings.
- `ordering`: This parameter is a string that controls the fill-reducing column preordering used. By default it is initialized as `None`, if this option is chosen, the solve functions use the appropriate default ordering (AMD for Cholesky and COLAMD for Left LU).
 - `'none'`: no column ordering; factorize A as-is.
 - `'colamd'`: COLAMD ordering
 - `'amd'`: AMD ordering

10.2 SPEX Python functions

10.2.1 `lu_backslash`: solves a linear system via exact LU factorization

The `lu_backslash` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times 1}$ and $\mathbf{b} \in \mathbb{R}^{n \times 1}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision.

The LU function expects as input a `scipy` sparse matrix A and a right hand side vector $\mathbf{b}$. Optionally, `option` object can be passed in. Currently, there are 2 ways to use this function outlined below.

- `x=SPEX.lu_backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution $\mathbf{x}$ is returned as a `numpy` double array.
- `x=SPEX.lu_backslash(A,b,options)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` object.

If the result $\mathbf{x}$ is held as a `numpy` double array, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution $\mathbf{x}$ may also be returned as a list of strings; See Section 10.1 for details.

10.2.2 `cholesky_backslash`: solves a linear system via exact Cholesky factorization

The `cholesky_backslash` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times 1}$ and $\mathbf{b} \in \mathbb{R}^{n \times 1}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision. Note that A must be symmetric positive definite.

The Cholesky function expects as input a `scipy` sparse matrix A and a right hand side vector $\mathbf{b}$. Optionally, `option` object can be passed in. Currently, there are 2 ways to use this function outlined below.

- `x=SPEX.cholesky_backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution $\mathbf{x}$ is returned as a `numpy` double array.
- `x=SPEX.cholesky_backslash(A,b,options)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` object.

If the result $\mathbf{x}$ is held as a `numpy` double array, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution $\mathbf{x}$ may also be returned as a list of strings; See Section 10.1 for details.

FIXME: where is ldl for Python?

10.2.3 `backslash`: solves a linear system via exact factorization

The `backslash` function solves the linear system $A\mathbf{x} = \mathbf{b}$ where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x} \in \mathbb{R}^{n \times 1}$ and $\mathbf{b} \in \mathbb{R}^{n \times 1}$. The final solution vector(s) obtained via this function are exact prior to their conversion to double precision. Note that A must be symmetric positive definite.

The Backslash function expects as input a `scipy` sparse matrix A and a right hand side vector $\mathbf{b}$. Optionally, `option` object can be passed in. If A is numerically symmetric, it attempts a Cholesky factorization. If the Cholesky fails or if the matrix is not numerically symmetric it performs an LU factorization. Currently, there are 2 ways to use this function outlined below.

- `x=SPEX.backslash(A,b)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using default settings. The solution `x` is returned as a `numpy` double array.
- `x=SPEX.backslash(A,b,options)` returns the solution to $A\mathbf{x} = \mathbf{b}$ using non-default settings from the `option` object.

If the result `x` is held as a `numpy` double array, in conventional floating-point representation (`double`), it is guaranteed to be exact only if the exact solution can be held in `double` without modification.

The solution `x` may also be returned as a list of strings; See Section [10.1](#) for details.

10.3 Python demo: example of how to use SPEX in Python

There is a file that provides a demo of the SPEX library in Python `demo.py`. It shows the usage for an exact solution as well as tuning the parameters.

- [1] P. R. AMESTOY, T. A. DAVIS, AND I. S. DUFF, *An approximate minimum degree ordering algorithm*, SIAM Journal on Matrix Analysis and Applications, 17 (1996), pp. 886–905.
- [2] ———, *Algorithm 837: AMD, an approximate minimum degree ordering algorithm*, ACM Transactions on Mathematical Software (TOMS), 30 (2004), pp. 381–388.
- [3] T. A. DAVIS, J. R. GILBERT, S. I. LARIMORE, AND E. G. NG, *Algorithm 836: COLAMD, a column approximate minimum degree ordering algorithm*, ACM Transactions on Mathematical Software (TOMS), 30 (2004), pp. 377–380.
- [4] ———, *A column approximate minimum degree ordering algorithm*, ACM Transactions on Mathematical Software (TOMS), 30 (2004), pp. 353–376.
- [5] L. FOUSSE, G. HANROT, V. LEFÈVRE, P. PÉLISSIER, AND P. ZIMMERMANN, *MPFR: a multiple-precision binary floating-point library with correct rounding*, ACM Transactions on Mathematical Software (TOMS), 33 (2007), p. 13.
- [6] T. GRANLUND ET AL., *GNU MP 6.0 Multiple Precision Arithmetic Library*, Samurai Media Limited, 2015.
- [7] C. LOURENCO, A. R. ESCOBEDO, E. MORENO-CENTENO, AND T. A. DAVIS, *Exact solution of sparse linear systems via left-looking roundoff-error-free LU factorization in time proportional to arithmetic work*, SIAM Journal on Matrix Analysis and Applications, 40 (2019), pp. 609–638.
- [8] C. J. LOURENCO AND E. MORENO-CENTENO, *Exactly solving sparse rational linear systems via roundoff-error-free cholesky factorizations*, SIAM Journal on Matrix Analysis and Applications, 43 (2022), pp. 439–463.

FIXME: add all of our SPEX papers to the list of references